

# **COMP5349– Cloud Computing**

## **Week 7: Elasticity and High Availability**

---

Dr. Ying Zhou

The University of Sydney

# Table of Contents

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

## WARNING

This material has been reproduced and communicated to you by or on behalf of the **University of Sydney** pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice

- 00 Federated User in AWS Academy Lab
- 01 Scalability and Elasticity
- 02 Scaling Compute Resource
- 03 Scaling Compute Resource
- 04 Elastic Load Balancing
- 05 Scaling Database Resource

# Federated User in AWS Academy Lab

# Session policy

---

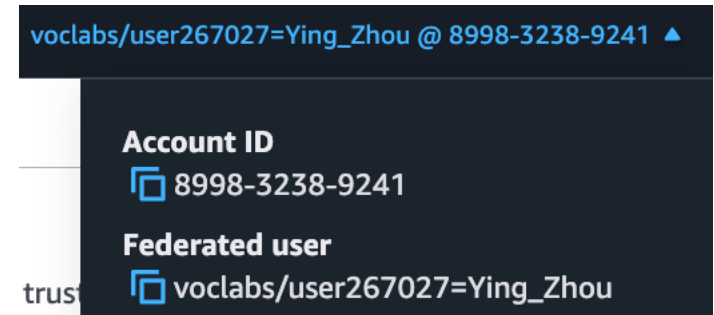
All federated users in AWS academy assume the role “voclabs”

The role is likely defined under a single or a few AWS accounts

Each lab has different sets of permissions (policies) associated with this role

At any time, different users can attempt different labs

Policies are dynamically associated with this role, and multiple policies can be associated with it at the same time



# Session policy

---

- Session policies are advanced policies that you pass as a **parameter** when you programmatically create a temporary session for a role or federated user.
- When AWS academy backend calls AssumeRole for a federated user, it includes a Policy parameter that limit the session's permissions

```
AssumeRole(  
  RoleArn="arn:aws:sts::899832389241:role/voclabs/",  
  RoleSessionName='131-5457484-8425041',  
  Policy='{  
    "Version": "2012-10-17",  
    "Statement": [  
      {  
        "Effect": "Allow",  
        "Action": ["s3:ListBucket", "ec2:DescribeInstances"],  
        "Resource": "*"   
      }  
    ]  
  }'  
)
```

# Lab Session vs. AWS Session

Start AWS session



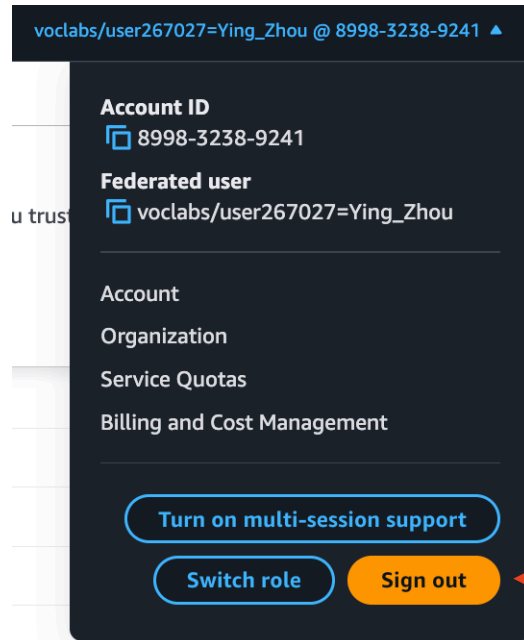
Start lab session

End lab session

01:32

▶ Start Lab

■ End Lab



When you end the lab session, all lab resources will be terminated, your federated user's permission will be revoked

But, your AWS session could still be alive and you can navigate the AWS console, but you can't do useful things now

When you end your AWS session before your lab session expires, your created resources are still there and you can start AWS session again by clicking the AWS button

End AWS session

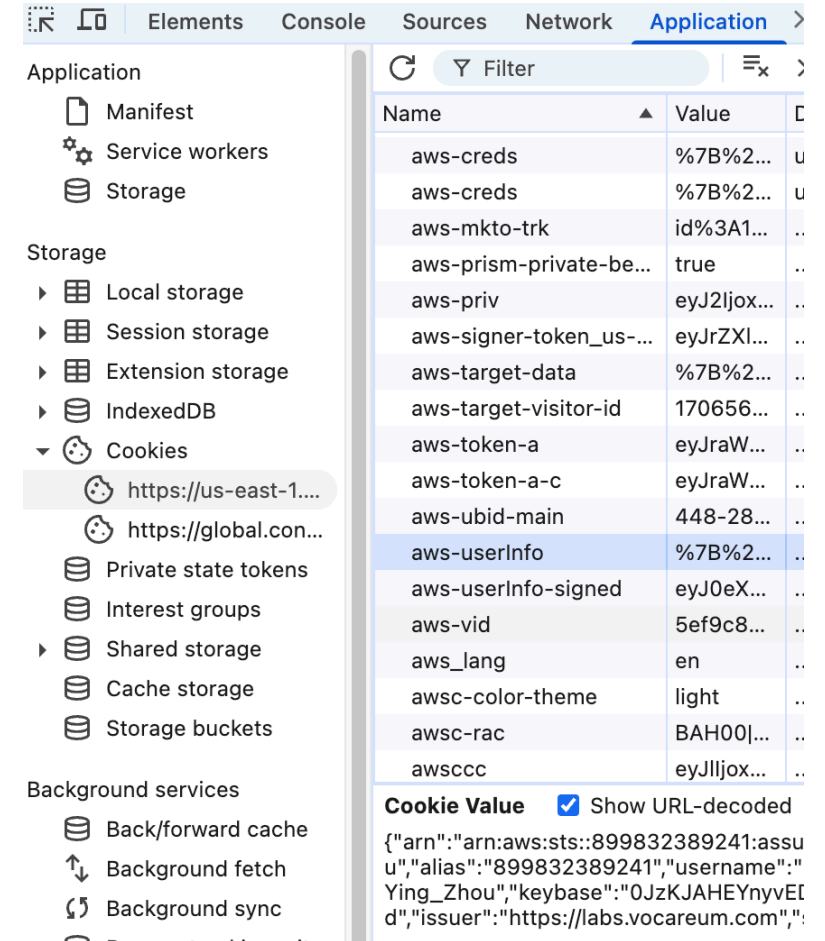
# What does start lab do?

---

- We are already authenticated by AWS academy
- When we click the “start lab” button
  - It calls AWS **STS (AssumeRole)** on a preconfigured IAM role — voclabs.
    - It passes a **custom session name** like user267027=Ying\_Zhou to track the authenticated user.
    - It may also attach a **session policy** to restrict what this session can do (lab-specific permissions).
  - AWS STS responds with temporary credentials
    - SessionToken, SecretAccessKey, etc
  - AWS Academy uses the information to generate a federated user sign-in URL
    - <https://signin.aws.amazon.com/federation?Action=login&SigninToken=<token>&Destination=https://console.aws.amazon.com/>
    - This is the url you will send when clicking the green AWS button
  - AWS Academy also uses a blueprint (CloudFormation template) to warm up the lab session by creating resources like a LabVPC, an EC2 instance

# When we click the green AWS session

- This URL grants us **browser-based access** to the AWS Console using the temporary credentials.
- The URL is opened in a new tab to display the AWS console
- When the AWS console opens, AWS set a lot of cookies to keep track of the session
  - aws-creds –contains the STS credentials
  - aws-sessionID – tracks the console session
  - Others



The screenshot shows the 'Application' tab in a web browser's developer tools. The left sidebar lists categories: Application, Storage, and Background services. Under 'Storage', 'Cookies' is expanded, showing two cookies for 'https://us-east-1...' and 'https://global.con...'. The 'aws-userInfo' cookie is selected. The right pane displays a table of cookies with columns 'Name', 'Value', and 'Domain'. The 'aws-userInfo' cookie is highlighted, showing its value as a JSON object containing STS credentials and user information.

Name	Value	Domain
aws-creds	%7B%2...	u
aws-creds	%7B%2...	u
aws-mkto-trk	id%3A1...	..
aws-prism-private-be...	true	..
aws-priv	eyJ2ljox...	..
aws-signer-token_us-...	eyJrZXl...	..
aws-target-data	%7B%2...	..
aws-target-visitor-id	170656...	..
aws-token-a	eyJraW...	..
aws-token-a-c	eyJraW...	..
aws-ubid-main	448-28...	..
aws-userInfo	%7B%2...	..
aws-userInfo-signed	eyJ0eX...	..
aws-vid	5ef9c8...	..
aws_lang	en	..
awsc-color-theme	light	..
awsc-rac	BAH00j...	..
awsc	eyJlljox...	..

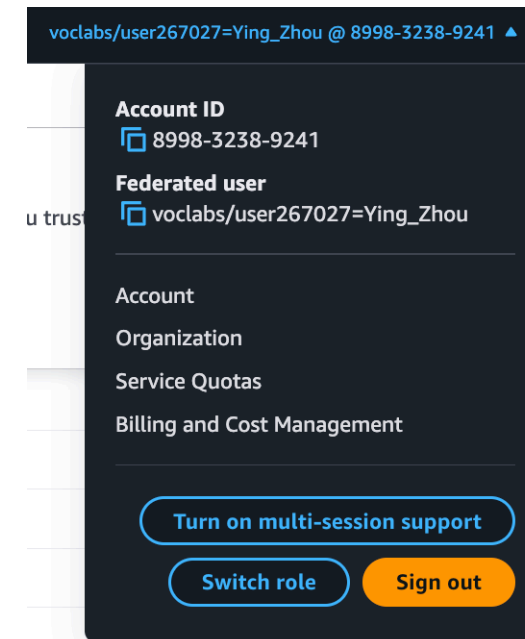
**Cookie Value** ☒ Show URL-decoded

```
{ "arn": "arn:aws:sts::899832389241:assu  
u", "alias": "899832389241", "username": "  
Ying_Zhou", "keybase": "0JzKJAHEYnyvEI  
d", "issuer": "https://labs.vocareum.com", "
```



# When we sign out from AWS

- All cookies will be cleared
- But, as long as the lab session is still on, we can click the green AWS button to start a new session
- The button is associated with the same federated user sign-in URL
- When we access the AWS console within the same lab session, we are recognized by AWS as the same user



# What does end lab do

AWS 

01:32

▶ Start Lab

■ End Lab

- The lab session is managed by AWS academy
- When we click end lab
  - All resources created during the lab are automatically cleared
  - The session permissions (session policy) are revoked or disabled
  - The AWS button is disabled
- If the AWS Console is still open in a browser tab, you may still **see the UI**, but most actions will fail due to permission loss.
- AWS Academy **cannot clear the AWS Console cookies** due to the browser's same-origin policy.
- You are required to explicitly logout the current AWS session to visit the console of a new lab

# Session expiration

---

- Both AWS and lab session expire after inactivity or a set time.
- When AWS session expires
  - The effect is the same as explicitly logging out – you lose access to the resources
  - If your **lab session is still valid**, you may be able to launch the console again and resume your work.
- When the lab session expires
  - The effect is the same as clicking the end lab button
  - All resources are cleared automatically and the permissions are revoked or disabled
  - Your temporary credential becomes invalid
  - You must click “**Start Lab**” again to get new credentials and restart the environment.

# What does click start lab again do?

---

- Core action
  - This resets the timer of a current lab session, giving you additional time to work on the same lab
  - Each lab session has a fixed duration (e.g., 2 hours), which is extended each time you click “Start Lab”
- What does not change
  - Your AWS Console credentials (temporary STS tokens) do not always refresh unless they are close to expiration.
  - Your assumed role and session information remain the same
  - Resources created during the lab session remain intact

# Scalability and Elasticity

# What is scalability

- **Scalability** is the property of a system to handle a *growing* amount of work.

<https://en.wikipedia.org/wiki/Scalability>

Usually judged by how easy it is to add more resources to handle such changes

An example of a system not scalable



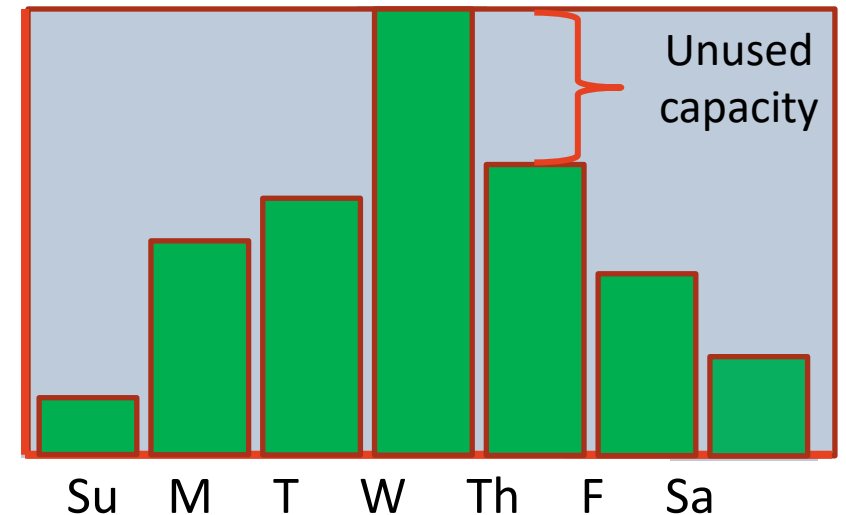
# What is elasticity?

---

- Supports fine grained capacity expansion and *contraction*.
  - Scalability in cloud setting
- Scalability in traditional setting is mostly about capacity expansion

## Examples:

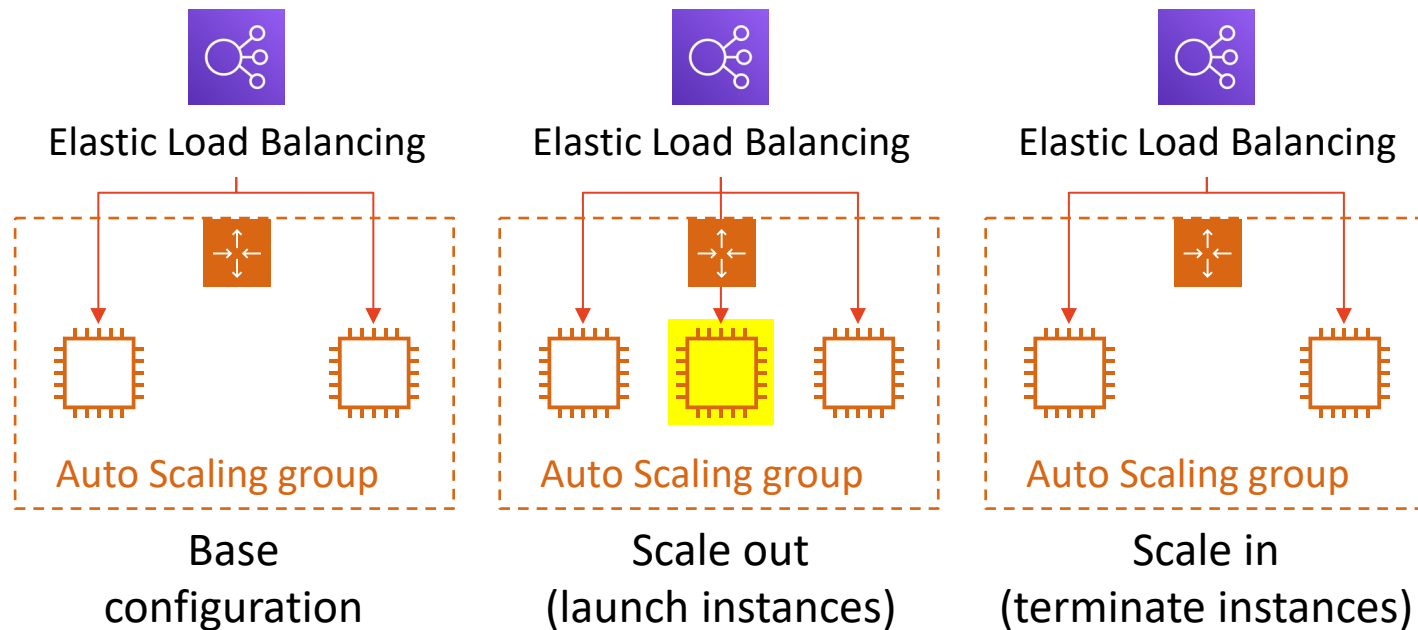
- Increasing the number of web servers when traffic spikes
- Lowering write capacity on your database when traffic goes down
- Handling the day-to-day fluctuation of demand throughout your architecture




# Scaling a computer system

- Is implemented by adjusting the amounts of resources based on the workload

## Horizontal scaling



## Vertical scaling





# What are required to achieve elasticity in cloud

---

- The ability to monitor the workload and make scaling decisions accordingly
  - The cloud platform provides the mechanism, e.g. **CloudWatch**, **AWS Auto Scaling**
  - The users need to configure them properly
- The ability to automatically start/stop new resources with minimal latency
  - The cloud platform provides the mechanisms, e.g. EC2 instances can be created from AMI with very small latency
  - The users need to configure them properly
- The resource change should be recognized by the running system
  - Cloud platform provides the mechanisms, such as AWS Elastic load balancing that can direct requests to newly added resources
  - The users need to ensure that the functionalities of the application are not affected by adding or removing resources



# Resource Monitoring

# Amazon CloudWatch

---



Amazon  
CloudWatch



- Monitors –
  - AWS resources
  - Applications that run on AWS
- Collects and tracks –
  - Standard metrics
  - Custom metrics
- Alarms –
  - Send notifications to an Amazon SNS topic
  - Perform Amazon EC2 Auto Scaling or Amazon EC2 actions
- Events –
  - Define rules to match changes in AWS environment and route these events to one or more target functions or streams for processing

# CloudWatch alarms

- Create alarms based on –
  - Static threshold
  - Anomaly detection
  - Metric math expression
- Specify –
  - Namespace
  - Metric
  - Statistic
  - Period
  - Conditions
  - Additional configuration
  - Actions

Statistic

Q Average



Period

5 minutes



## Conditions

Threshold type



Static

Use a value as a threshold



Anomaly detection

Use a band as a threshold

Whenever CPUUtilization is...

Define the alarm condition



Greater

> threshold



Greater/Equal

>= threshold



Lower/Equal

<= threshold



Lower

< threshold

than...

Define the threshold value

100




Must be a number

► Additional configuration

# Scaling Compute Resource

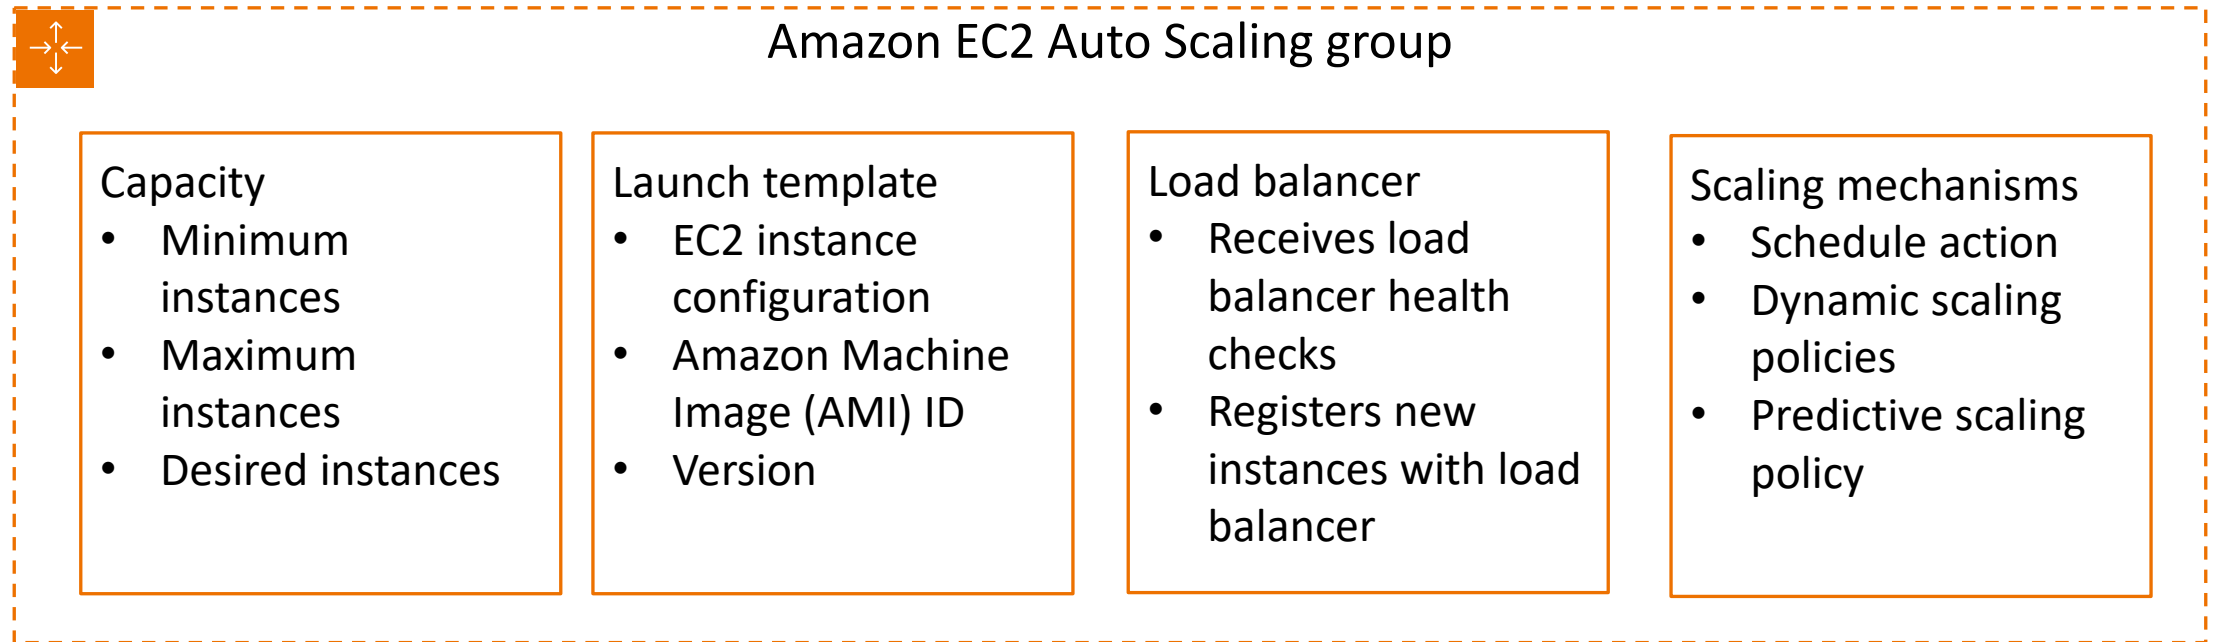
# Amazon EC2 Auto Scaling



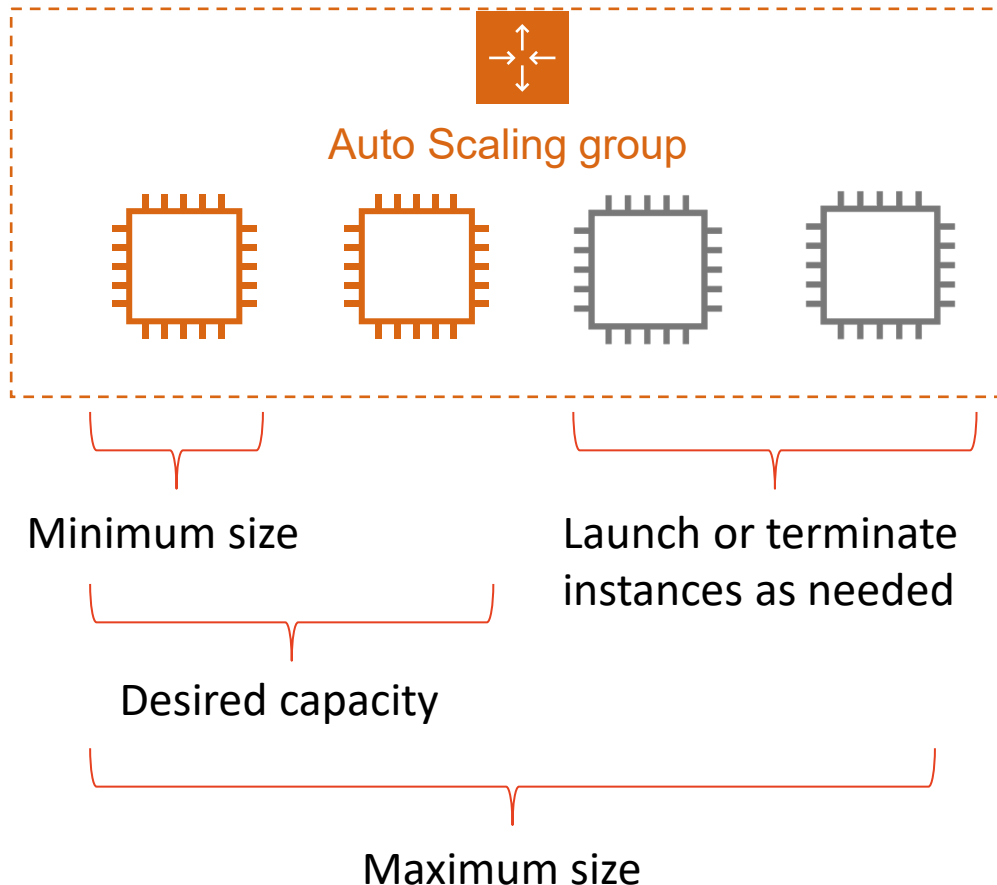
Amazon EC2 Auto  
Scaling

- Manages a logical collection of Amazon Elastic Compute Cloud (Amazon EC2) instances called an Amazon EC2 Auto Scaling group across Availability Zones
- Launches or retires EC2 instances configured by launch templates
- Resizes based on events from scaling policies, load balancer health check notifications, or schedule actions
- Integrates with Elastic Load Balancing (ELB) to send new instances registrations and receive health notifications
- Balances the number of instances across Availability Zones
- Is available free of charge

# Amazon EC2 Auto Scaling group components



# Auto Scaling Group - capacity

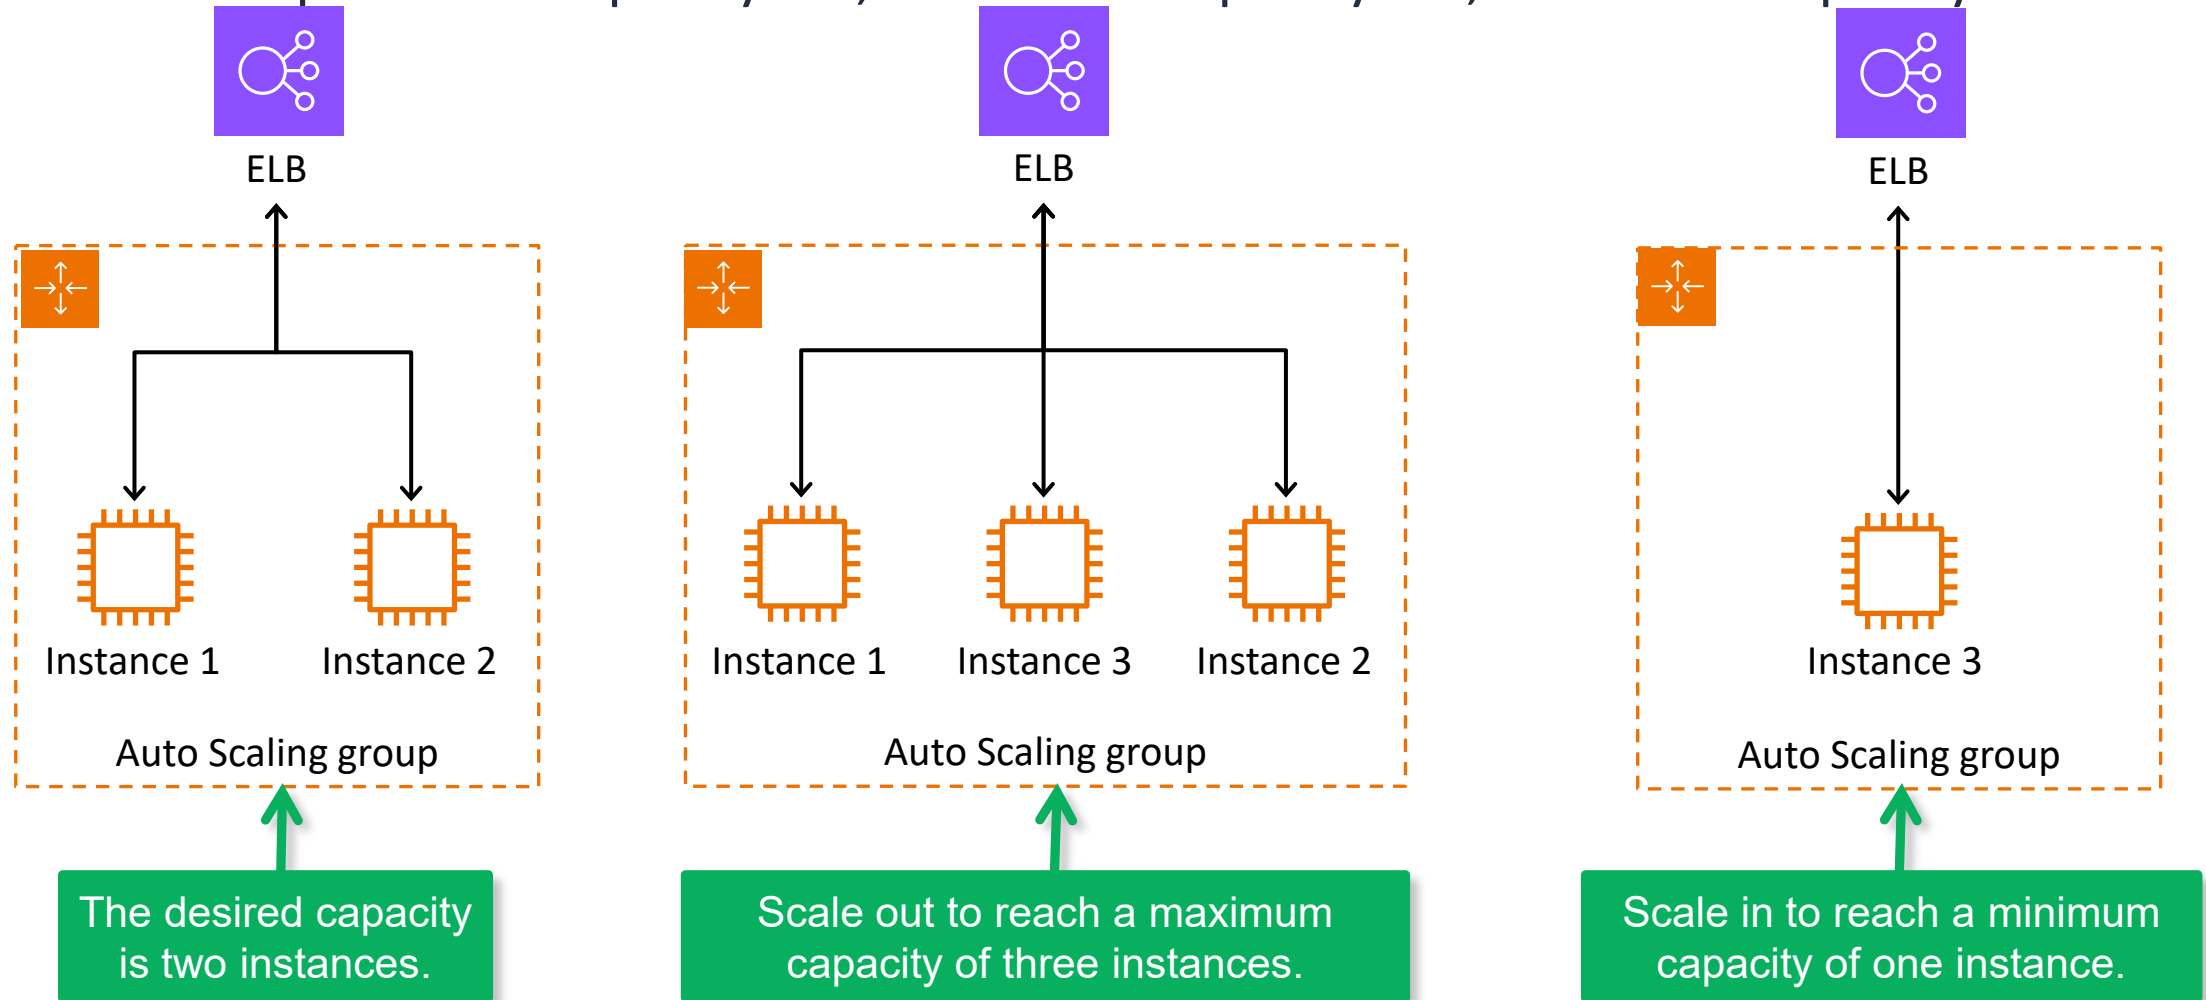


An **Auto Scaling group** is a collection of EC2 instances that are treated as a logical grouping for the purposes of automatic scaling and management.

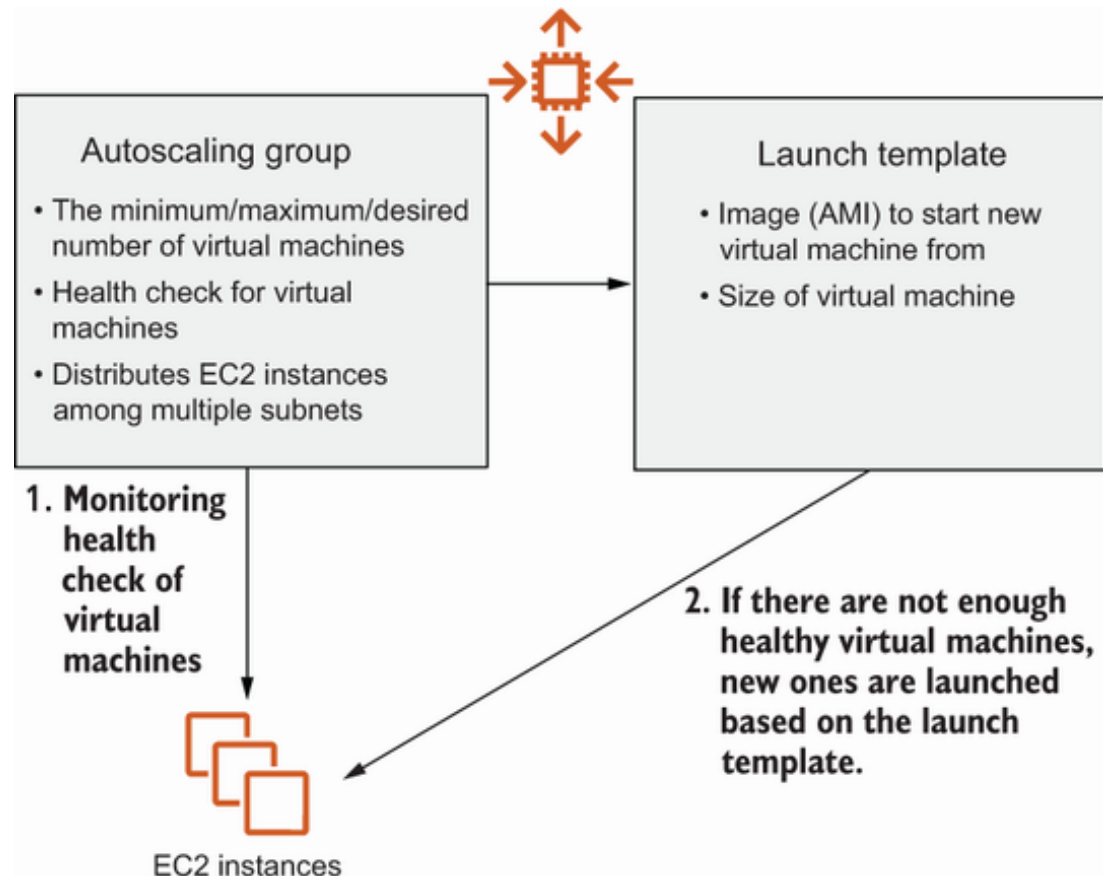


# Horizontal scaling with Amazon EC2 Auto Scaling groups

Group desired capacity = 2, maximum capacity = 3, minimum capacity = 1

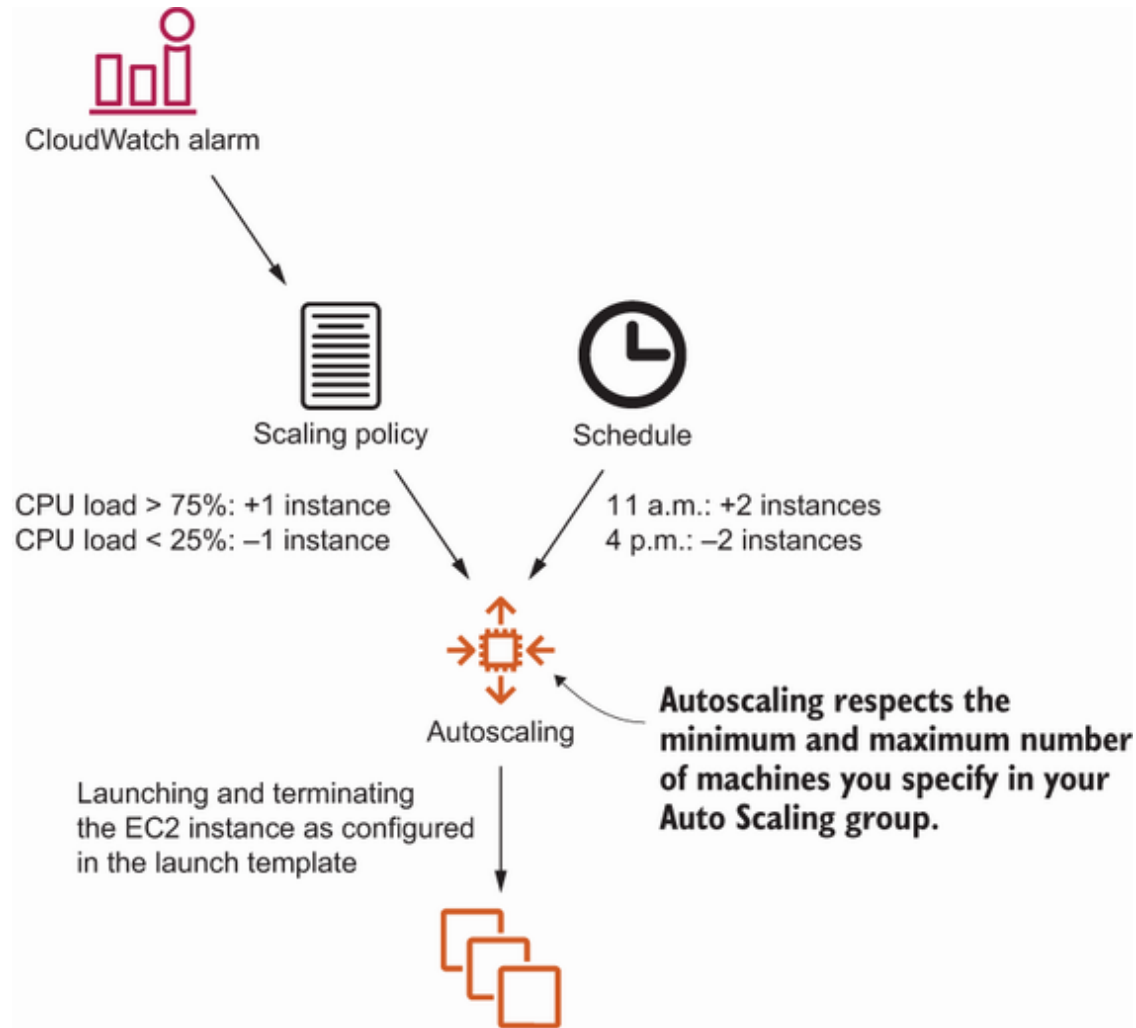


# Auto Scaling Group – Launch Template



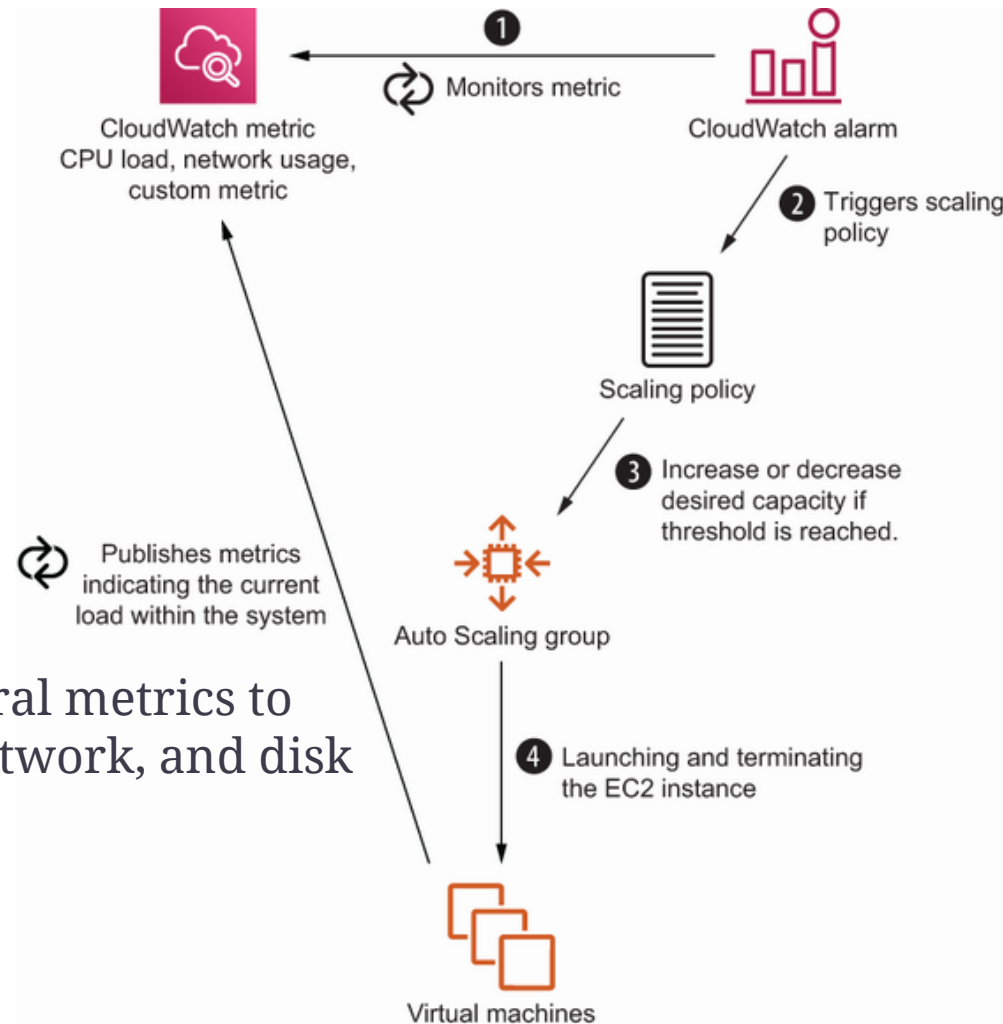
- A *launch template* contains all information needed to launch an EC2 instance
  - Instance type (size of virtual machine)
  - Image (AMI)
  - Security Group
  - User Data
- The subnet for the virtual machine isn't defined in the launch template but rather in the Auto Scaling group.

# Scheduling Mechanism



- Based on schedule
  - E.g. : Turning off your development and test instances at night
- Based on policy
  - Usually relying on CloudWatch alarm
  - Maintaining constant number
    - Always 2 instances
  - Target tracking
    - Maintaining an average CPU load of 60%
  - Other policies

# Scheduling policy based on CloudWatch Metrics



An EC2 instance publishes several metrics to CloudWatch by default: CPU, network, and disk utilization and others

# Elastic Load Balancing

# Elastic Load Balancing

---



Elastic Load  
Balancing

A **managed load balancing service** that distributes incoming application traffic across multiple EC2 instances, containers, IP addresses, and Lambda functions.


- Can be **external-facing or internal-facing**
- Each load balancer receives a **DNS name**
- Recognizes and responds to **unhealthy instances**

# Types of AWS load balancers




## Application Load Balancer

- Is used for HTTP and HTTPS traffic
- Operates at OSI layer 7, the application layer
- Is used for application architectures



## Network Load Balancer

- Is used for TLS offloading, UDP, and static IP addresses
- Operates at OSI layer 4, the transport layer
- Is used for millions of requests per second at ultra-low latency



## Gateway Load Balancer


- Is used for third-party virtual appliance fleet using GENEVE protocol
- Operates at OSI layer 3, the network layer
- Is used to improve security, compliance, and policy controls



## Classic Load Balancer

- Is used for previous generation EC2-Classic networks
- Operates at OSI layers 3 and 7, the transport and application layers
- Is used if upgrading to other load balancers is not feasible

# What does a load balancer do?



A load balancer is used to route incoming requests to multiple EC2 instances, containers or other targets

The server running on the EC2 instance must be stateless

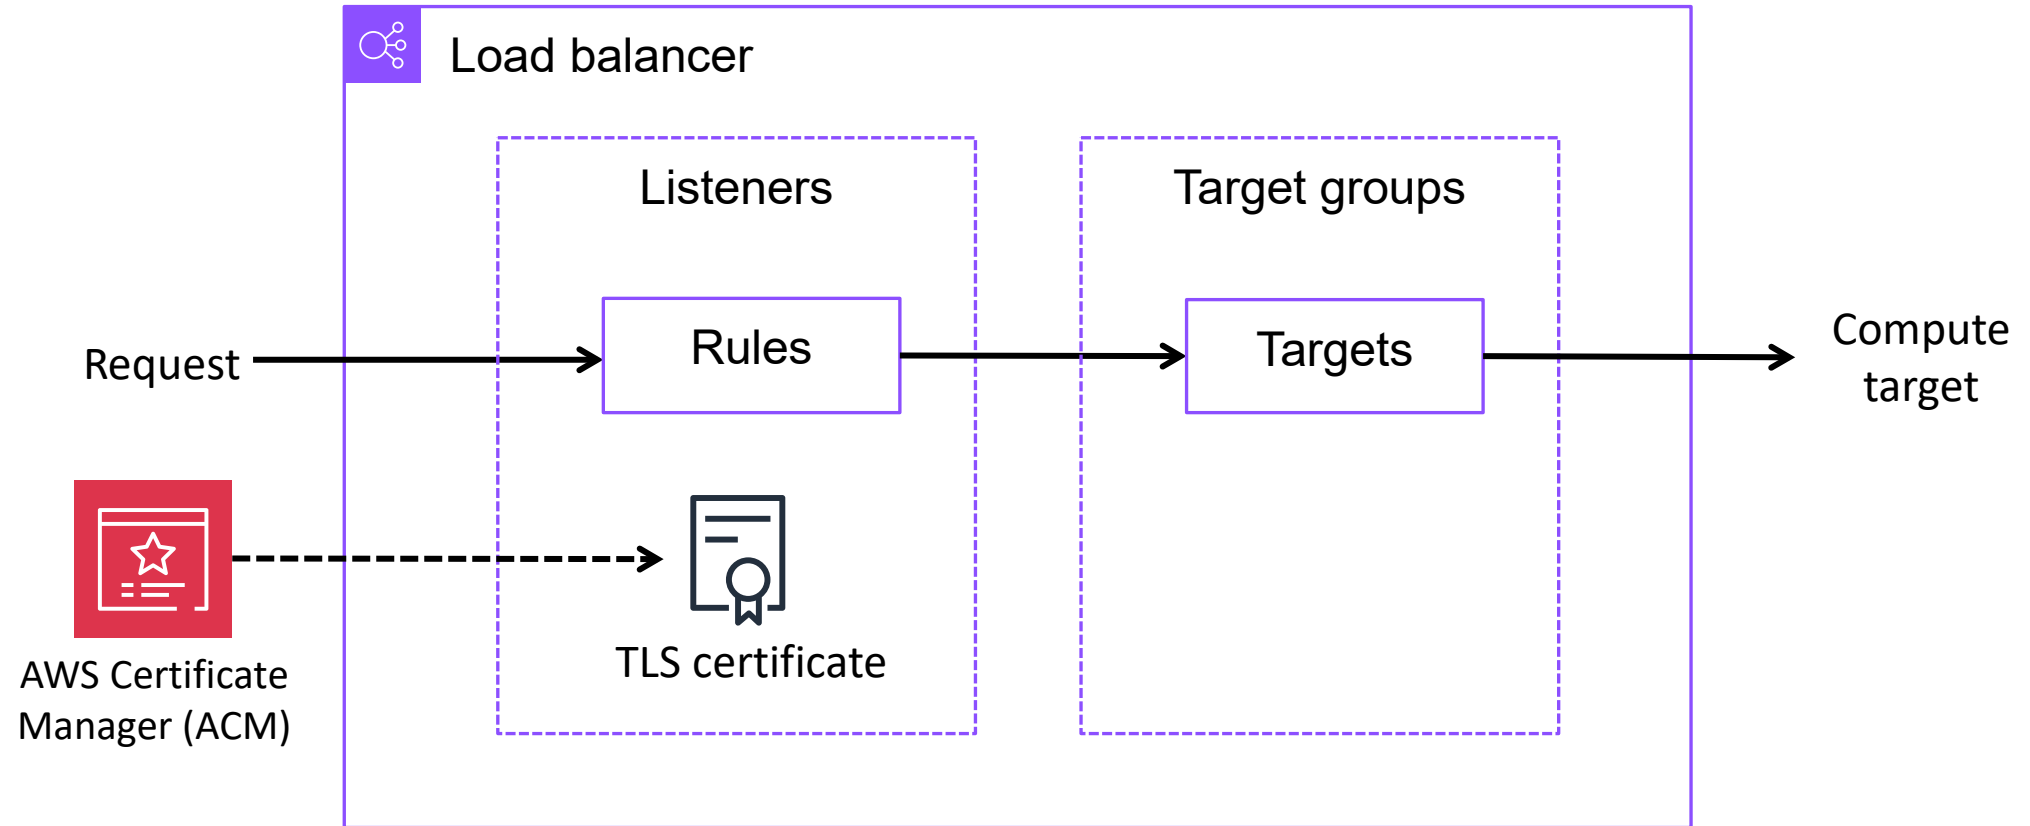
A stateless server stores any shared data remotely in a database or storage system.

When you create an ELB and enable it in one or more Availability Zones (AZs), AWS provisions *load balancer nodes* within those AZs. These nodes are the actual components that receive incoming traffic and distribute it to your backend instances.

The capacity can grow or shrink depends on the workloads

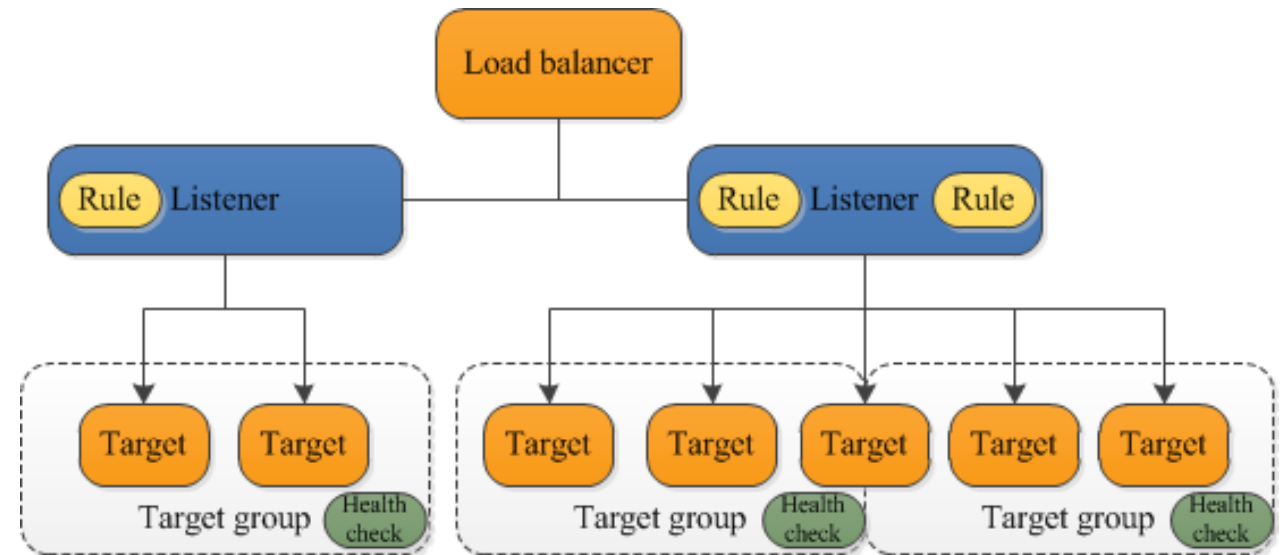


# Load balancer components



# ALB Components

- Listener
  - Defines the port and protocol which the load balancer must listen on
  - Each ALB must have at least one listener
  - Each ALB can have multiple listeners
  - Listener defines routing rules
- Target Group
  - Logical grouping of AWS resources as target for the traffic
    - E.g. An autoscaling group can be specified as the target group
  - Load balancer performs periodic health check on targets in the target groups. The request will only be routed to healthy target




# Listeners and rules

---

- Application load balancers are used to distribute HTTP and HTTPS requests
  - They support HTTP and HTTPS listeners
  - One ALB must have one listener
- Each listener defines routing rules for incoming requests
  - Rules consists of conditions and actions
  - When a request meets the condition of the rule, the action is taken
  - Typical actions
    - forwarding to a target group
    - redirect to a url
    - Responding back to the client

# An example listener

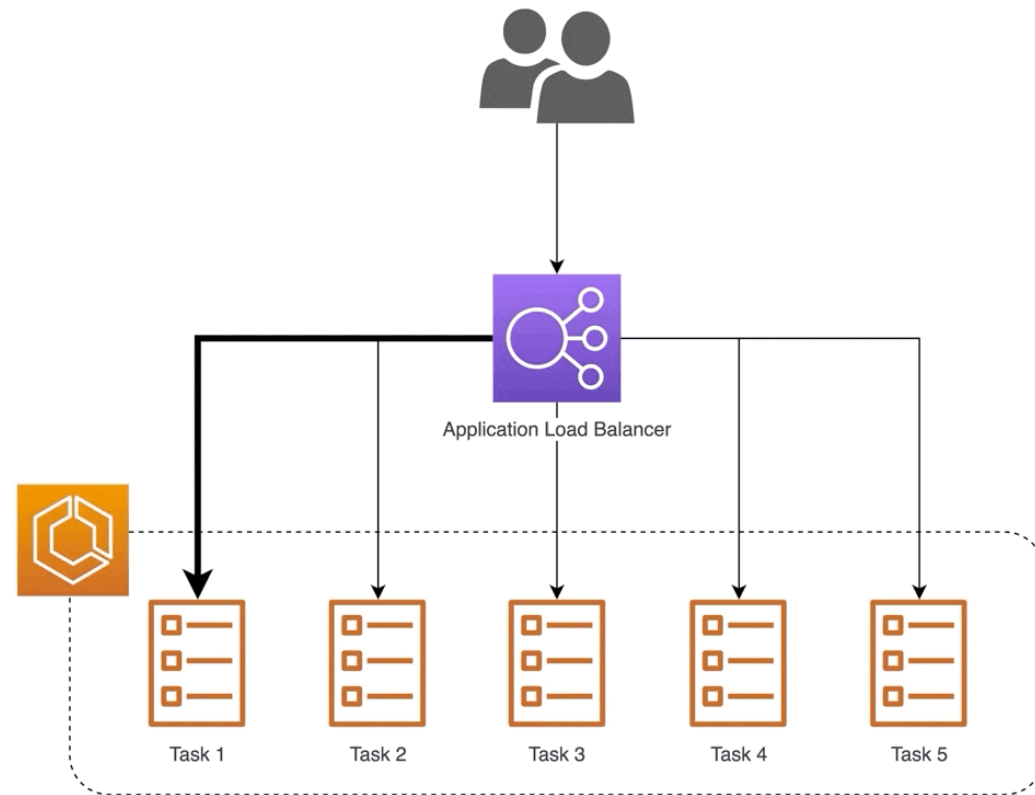


Priority	Conditions (If)	Actions (Then) <a href="#">↗</a>
Last (default)	<i>If no other rule applies</i>	<b>Forward to target group</b> <ul style="list-style-type: none"><li>• <a href="#">my-targets</a>: 1 (100%)</li><li>• Group-level stickiness: Off</li></ul>

Default rule

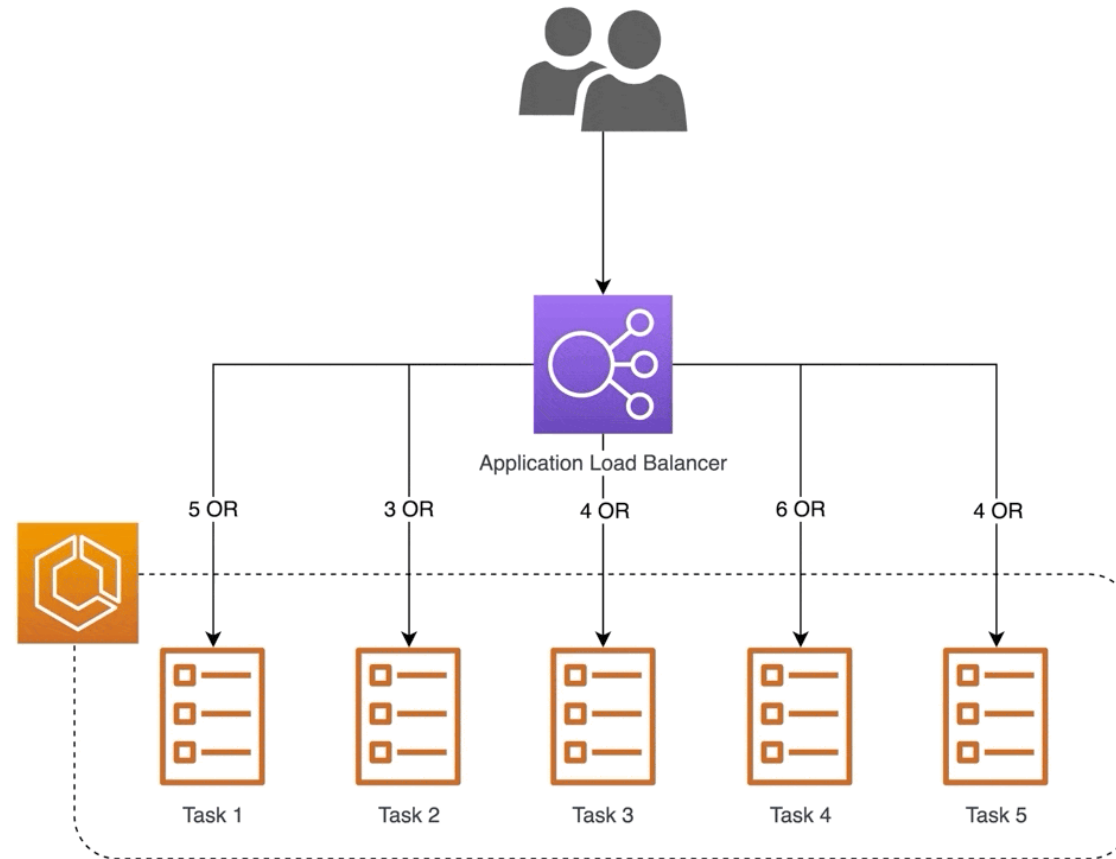
# Picking a target to forward the request

- All load balancers use some load balancing algorithms to distribute incoming requests among targets in a target group.
- ALB uses the following two algorithms
  - Default: Round-robin algorithm
    - Cycling through target
  - Configurable: Least Outstanding Request algorithm



# Least Outstanding Request algorithm

- LOR selects the target with the lowest number of requests waiting for responses.

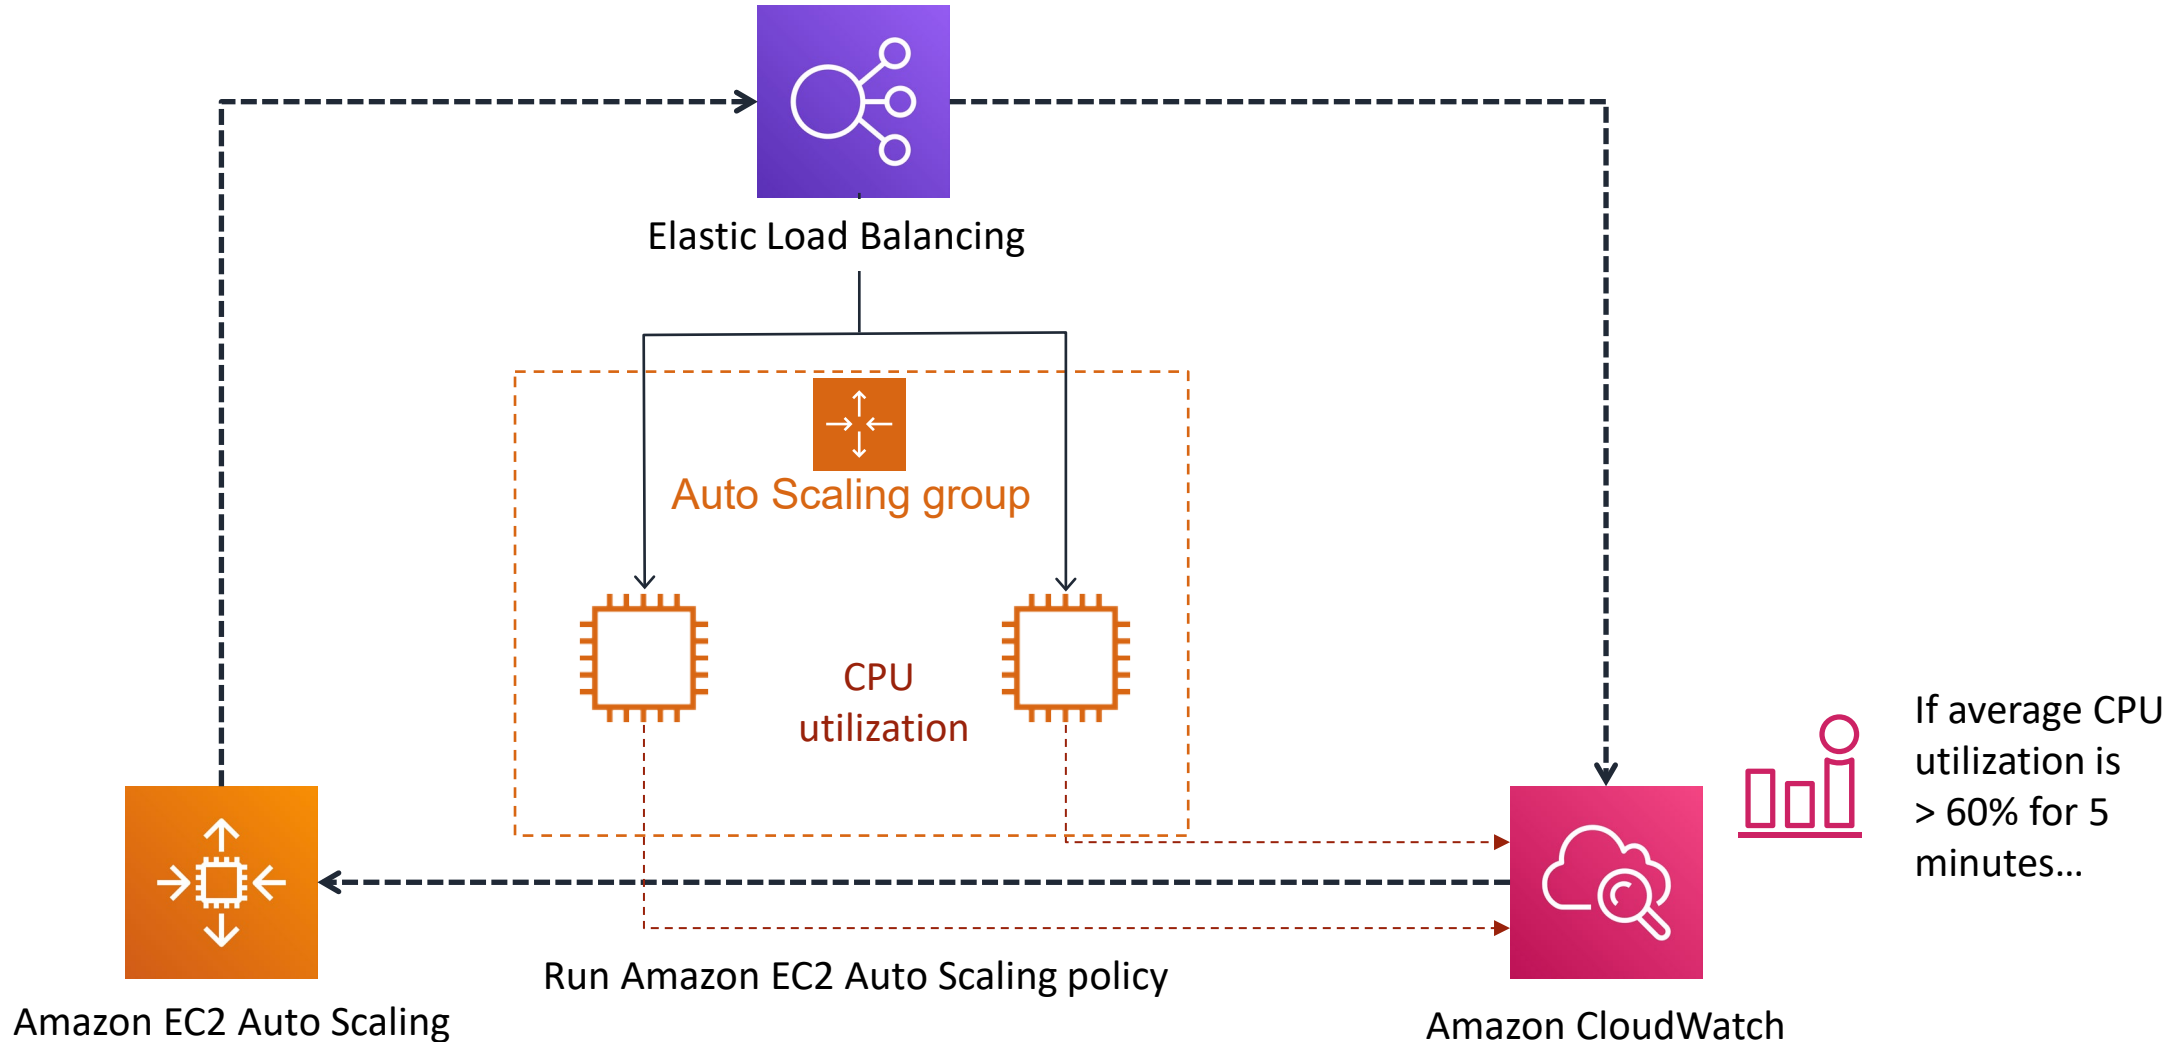


# Target group health check

---

- Application Load Balancer periodically sends requests to its registered targets to test their status. These tests are called *health checks*.
  - Requests are sent to healthy targets unless there is no healthy target
- The user needs to configure the health check settings
  - Protocol
  - Path
  - Timeout seconds
  - Interval
  - Etc,

# Dynamic scaling and load balancing





# ELB health check and ASG health check

---

- Auto Scaling Group always have EC2 level health check turned on
  - EC2 Auto Scaling checks that instances are running and monitors for underlying hardware or software issues that might impair an instance.
- It can be configured to use ELB health check as well
  - If the target group marks an instance as unhealthy, the ASG will terminate it and depending on the policy may start a new instance
  - Can detect application-level issue and replace instance accordingly

# Example Setting

---

- An **Auto Scaling Group (ASG)** managing **3 EC2 instances**.
- An **Application Load Balancer (ALB)** with a **Target Group** that receives traffic and routes it to instances in the ASG.
  - The **Target Group** does HTTP health checks every 30 seconds.
- The **ASG** is configured to use **ELB health checks**
  - it considers instances unhealthy if the Target Group marks them unhealthy

# Scenario: App Crash

---

- The app crashes, but the EC2 instance is still “up”
- Target group health check fails
- Since the ASG is using ELB health checks, it also:
  - Marks instance as unhealthy.
  - Terminates it and launches a replacement to keep 3 instances in the group

# Scenario: unresponsive instance

---


- EC2 health check fail
  - Detected by CloudWatch and reported to ASG
- ASG will terminate the instance and start a new one
  - This could eventually be detected by the target group health check
  - The CloudWatch detection of system failure is faster and ASG can react immediately

# Scaling Database Resource

# Vertical scaling with Amazon RDS: Push-button scaling

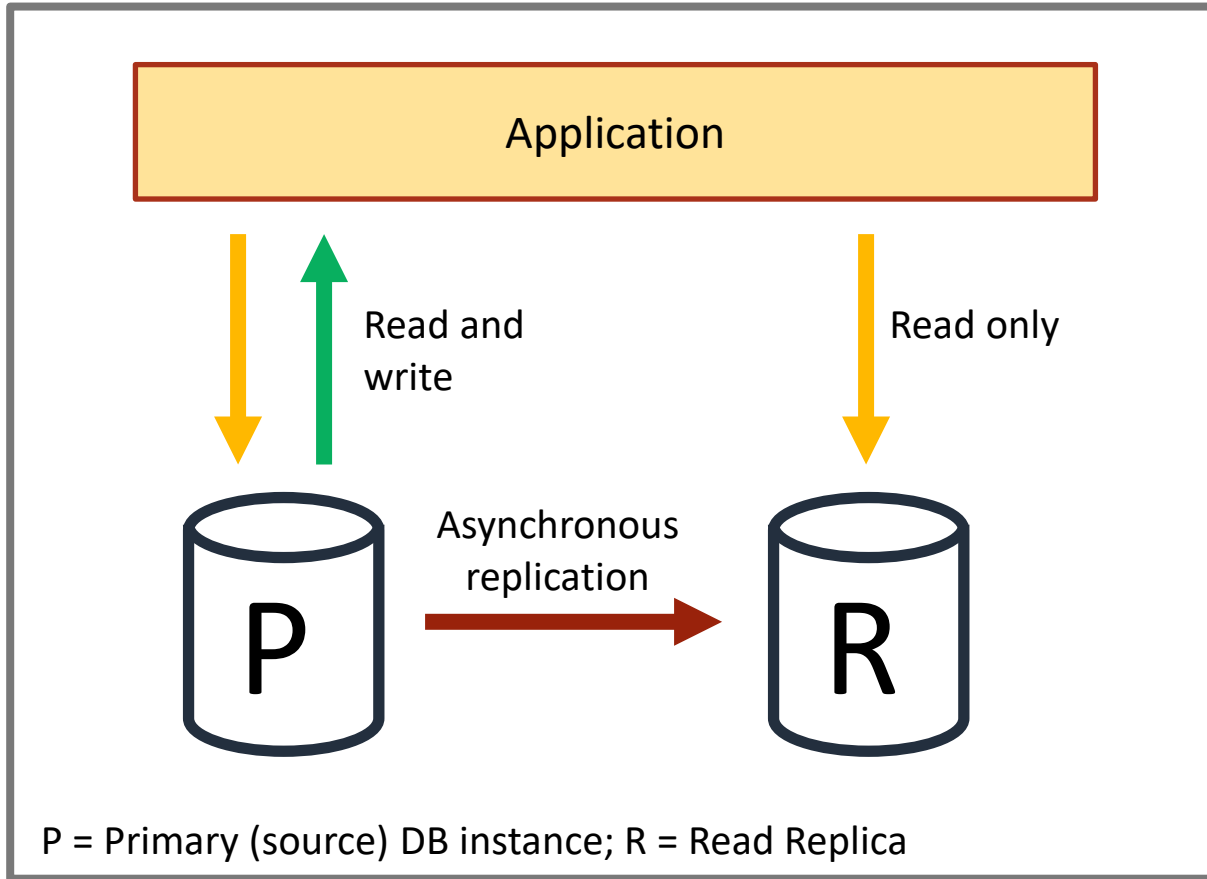
---

- Scale DB instances **vertically** up or down
- From **micro to 24xlarge** and everything in between
- Scale vertically with **minimal downtime**



# Horizontal scaling with Amazon RDS: Read replicas

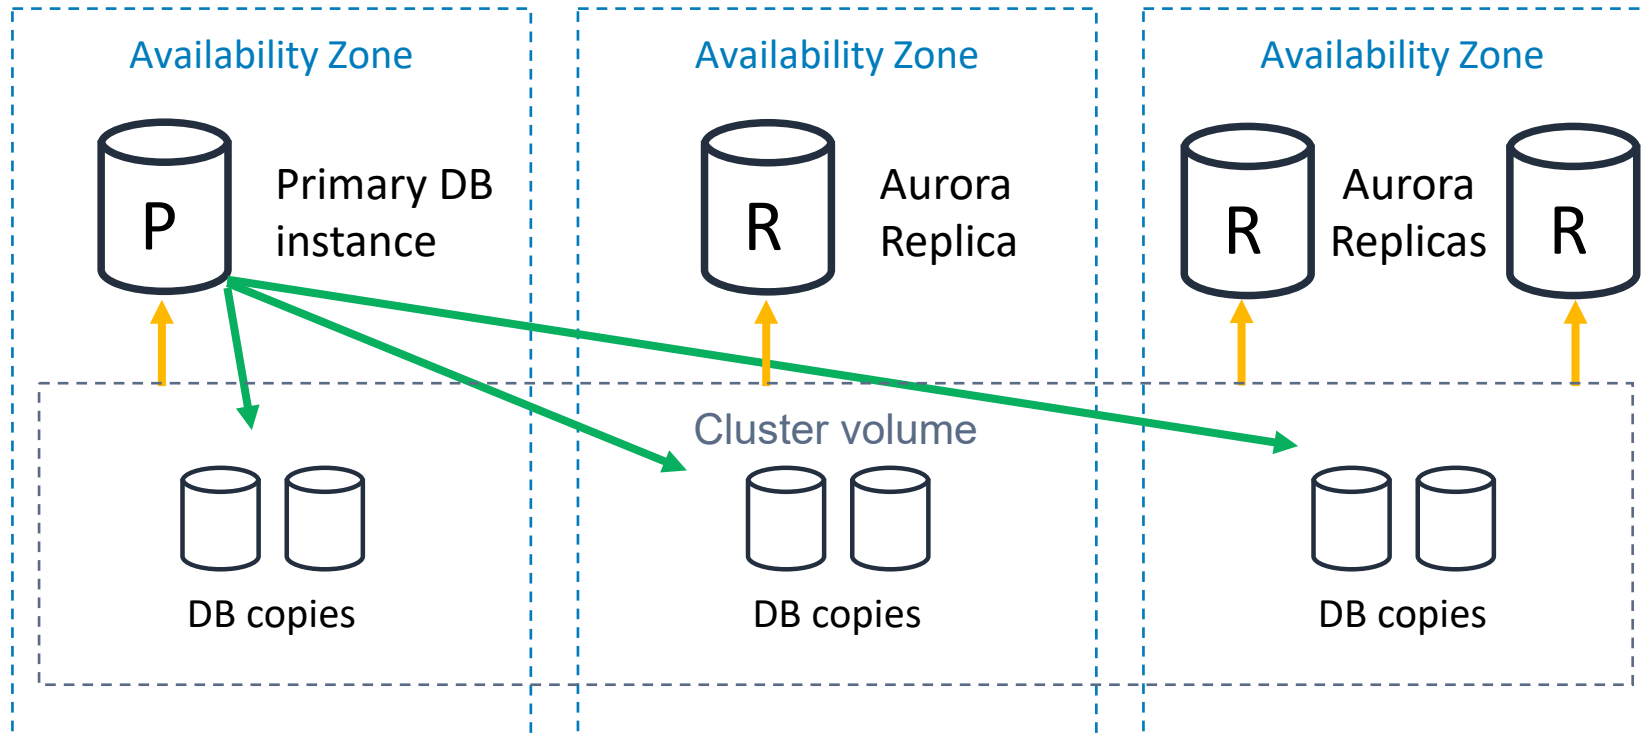
---




- Horizontally scale for **read-heavy** workloads
- Up to **five read replicas** and up to **15 Aurora replicas**
- Replication is **asynchronous**
- Available for Amazon RDS for MySQL, MariaDB, PostgreSQL, and Oracle

# Scaling with Amazon Aurora

Each Aurora DB cluster can have up to 15 Aurora replicas







THE UNIVERSITY OF  
**SYDNEY**