

Web Application Development

COMP4347

COMP5347

Node.js

MVC Architecture

Introduction to MongoDB

Week 6


Semester 1, 2025

Dr. Mehdi Nobakht

School of Computer Science



THE UNIVERSITY OF
SYDNEY



Outline

- Implementing MVC basics
 - Application Folder Structure
 - CommonJS modules
 - Controller and Routers
- Session Management
- Database Layer
 - Introduction to NoSQL
 - Introduction to MongoDB

The Survey app (Week 5 lab)

```
var express = require('express')
var path = require('path')
```

Application scope variables

```
var app = express()
```

```
var products=['iphone 7', 'huawei p9', 'Pixel XL', 'Samsung S7']
var surveyresults = { fp:[0,0,0,0],mp:[0,0,0,0]}
```

```
app.use(express.json());
app.use(express.urlencoded({extended: true}));
app.use(express.static(path.join(__dirname, 'public')));
app.set('views', path.join(__dirname, 'views'));
```

```
app.get('/', function(req,res){
    res.render('survey.pug',{products:products})
});
```

```
app.post('/survey', function(req,res){
    console.log(req.body);
    cast vote = req.body;
    ...
    res.render('surveyresult.ejs', {products: products, surveyresults: surveyresults})
});
app.listen(3000, function () {
    console.log('survey app listening on port 3000!')
})
```

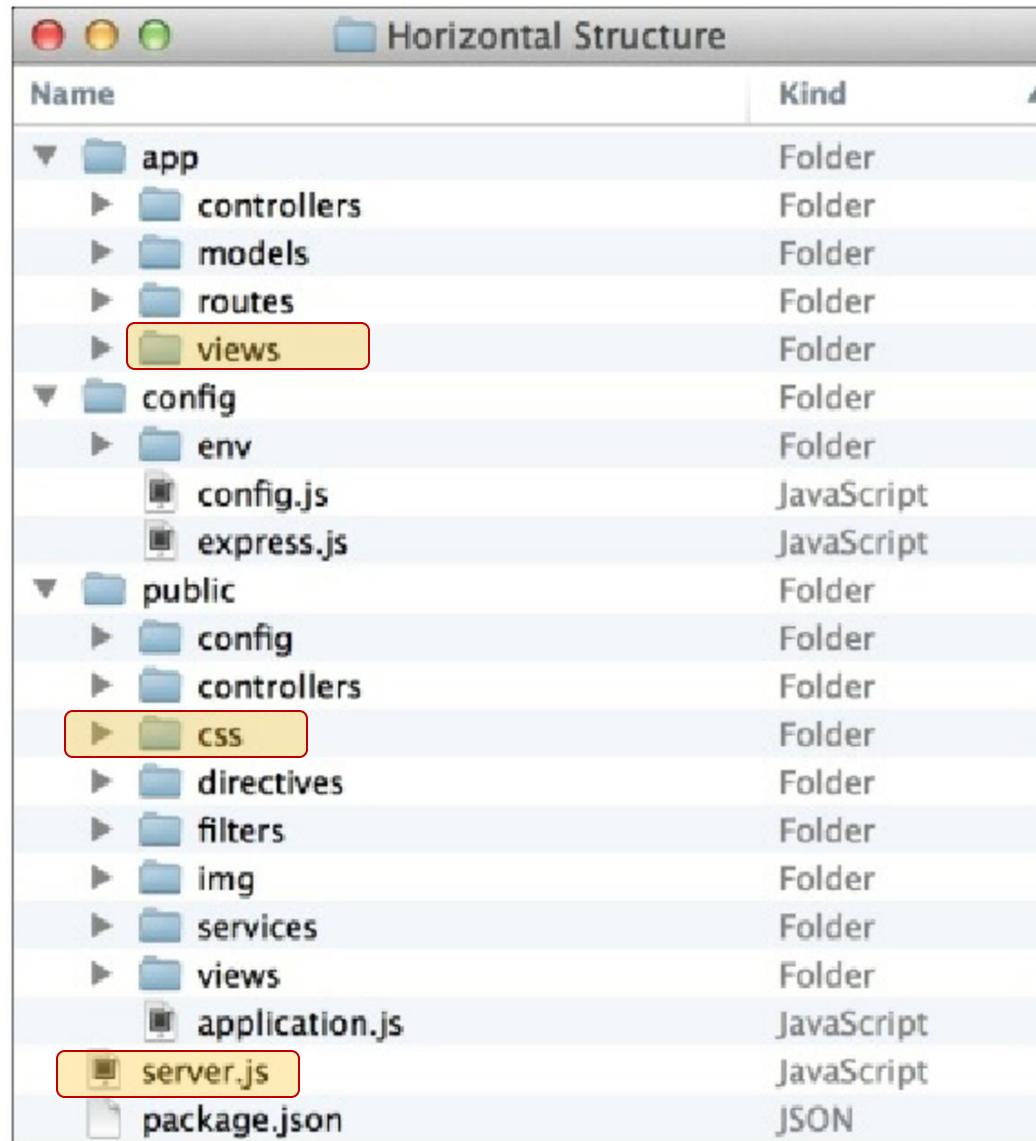
Request scope variables

Several Issues

```
✓ COMP5347
  > node_modules
  ✓ public/css
    # surveystyle.css
  ✓ views
    <> survey.ejs
    🍷 survey.pug
    <> surveyresult.ejs
  JS app.js
  {} package-lock.json
  {} package.json
```

- There is a *single* JS file with all settings and route methods
- Large application will have many route methods and some are related
 - Modular controller and route mappings
- Data sharing
 - Application scope variable
 - Request scope variable
 - Session scope variable
- We don't have separated model yet!

Application folder structure



Name	Kind
▼ app	Folder
▶ controllers	Folder
▶ models	Folder
▶ routes	Folder
▶ views	Folder
▼ config	Folder
▶ env	Folder
config.js	JavaScript
express.js	JavaScript
▼ public	Folder
▶ config	Folder
▶ controllers	Folder
▶ css	Folder
▶ directives	Folder
▶ filters	Folder
▶ img	Folder
▶ services	Folder
▶ views	Folder
application.js	JavaScript
server.js	JavaScript
package.json	JSON

CommonJS Module in Node.js

- A file-based module system to solve JavaScript single global namespace issue
 - Each file is its own module
- Three key components:
 - **requires()**: this method is used to load the module into the current code
 - Eg: `require(express)`
 - **exports**: this object is contained in each module and allows you to expose piece of your code when the module is loaded.
 - **module**: refers to the current module definition (metadata).

Writing our own module

hello.js

```
var message = 'Hello';  
  
module.exports.sayHello=function(){  
    console.log(message);  
}  
  
exports.sayBye=function(){  
    console.log("Bye")  
}
```

module.exports and **exports** are equivalent, both referring to the object exposed by the module

We can expose many methods by defining them as properties of the **module.exports** object.

Hello_client.js

```
var hello = require('./hello');  
hello.sayHello()  
hello.sayBye()
```

Calling **require(...)** in the client code would return the **module.exports** object. Our hello which has exposed two methods

Writing Controller(s)

survey.server.controller.js

```
var express = require('express')

module.exports.showForm = function(req,res){
  products = req.app.locals.products
  res.render('survey.pug',{products:products})
}

module.exports.showResult = function(req,res){
  console.log(req.body);
  gender = req.body.gender
  productidx = req.body.vote;
  products = req.app.locals.products;
  surveyresults = req.app.locals.surveyresults;
  if (gender == 0)
    surveyresults.mp[productidx]++;
  else
    surveyresults.fp[productidx]++;

  res.render('surveyresult.ejs ', {products: products,
    surveyresults: surveyresults})
}
```

This controller module exposes two methods: **showForm** is used for displaying the form; **showResult** is used for showing the results

The methods are not mapped to URL yet

req.app.locals is used to share application scope variables

Each request object has a reference to the current running express application: **req.app**

app.locals is used to store properties that are local variables within the application (application scope data)

Mapping Controller to URL

survey.server.routes.js

```
var express = require('express')
var controller = require('../controllers/survey.server.controller')
var router = express.Router()

router.get('/', controller.showForm)
router.post('/survey', controller.showResult)
module.exports = router
```

- **Express Router** middleware used to group the route handlers for particular part of a web application (e.g., user account functions)
- Router can be used to define the routes (HTTP request method, URL path/pattern, and callback function to handle that pattern)
- Using Router is similar to defining routes directly on the Express application object
 - Router supports application modularity

Mapping Controller to URL

```
var express = require('express')
```

```
var controller = require('../controllers/survey.server.controller')
```

```
var router = express.Router()
```

```
router.get('/', controller.showForm)
router.post('/survey', controller.showResult)
module.exports = router
```

survey.server.routes.js

```
var express = require('express');
```

```
var path = require('path')
```

```
var survey = require('../routes/survey.server.routes')
```

```
var app = express()
```

```
app.locals.products=['iphone 7', 'huawei p9', 'Pixel XL', 'Samsung S7']
app.locals.surveyresults = {
  fp:[0,0,0,0], mp:[0,0,0,0]
}
```

```
app.use(express.json()); % to parse json data from the request
app.use(express.urlencoded({ extended: true}));
app.set('views', path.join(__dirname, '/app/views'));
app.use(express.static(path.join(__dirname, 'public')));
```

```
app.use('/', survey)
```

```
app.listen(3000, function () {
  console.log('survey app listening on port 3000!')
```

```
})
```

server.js

Mapping Controller to URL

```
var express = require('express')
```

```
var controller = require('../controllers/survey.server.controller')
```

```
var router = express.Router()
```

```
router.get('/', controller.showForm)
```

```
router.post('/survey', controller.showResult)
```

```
module.exports = router
```

survey.server.routes.js

```
var express = require('express');
```

```
var path = require('path')
```

```
var survey = require('../routes/survey.server.routes')
```

Set the two application scope variables

```
var app = express()
```

```
app.locals.products=['iphone 7', 'huawei p9', 'Pixel XL', 'Samsung S7']
```

```
app.locals.surveyresults = {  
  fp:[0,0,0,0], mp:[0,0,0,0]  
}
```

```
app.use(express.json()); % to parse json data from the request
```

```
app.use(express.urlencoded({ extended: true}));
```

```
app.set('views', path.join(__dirname, '/app/views'));
```

```
app.use(express.static(path.join(__dirname, 'public')));
```

```
app.use('/', survey)
```

```
app.listen(3000, function () {
```

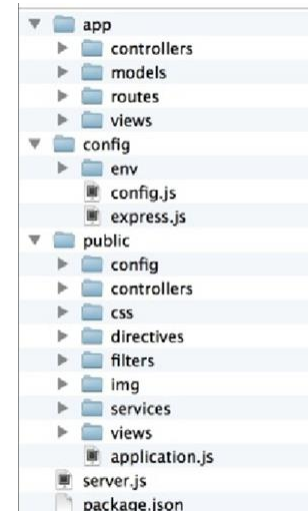
```
  console.log('survey app listening on port 3000!')
```

```
})
```

COMP5347 Web Application Development

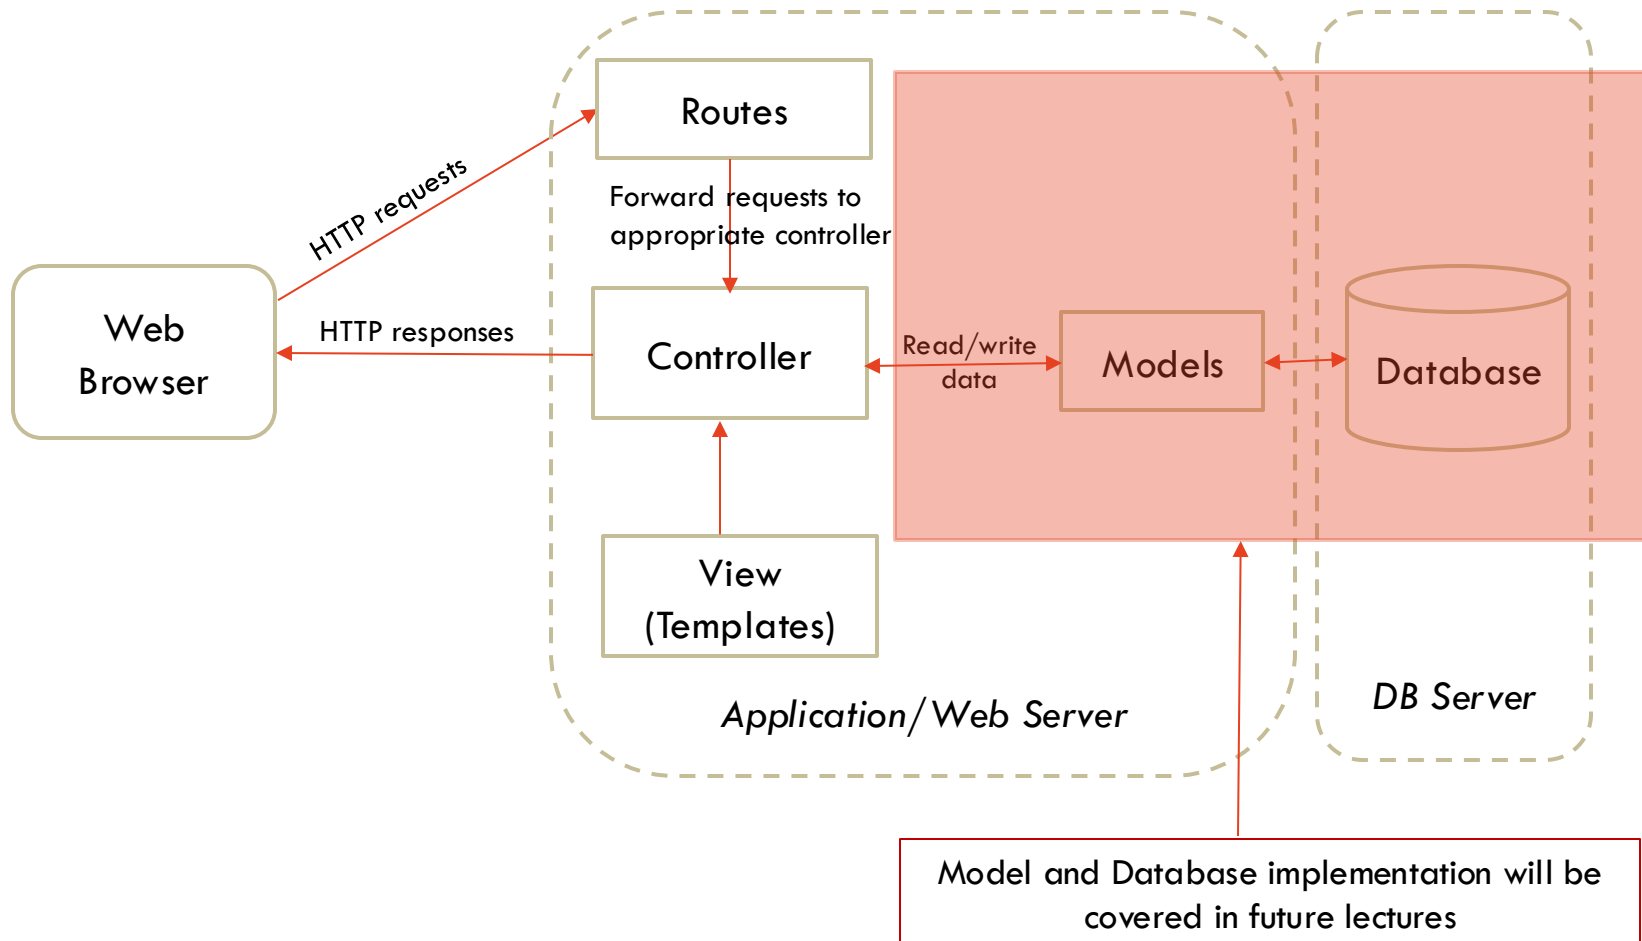
The University of Sydney

server.js



- Get request send to /survey will display the form
- Post request send to /survey/survey will show the result

Survey Web Application – MVC



Outline

- Implementing MVC basics
 - Application Folder Structure
 - CommonsJS modules
 - Controller and Routers
- **Session Management**
- Database Layer
 - Introduction to NoSQL
 - Introduction to MongoDB

Variable Scopes

- Application scope data is available through out the application
- Request scope data is available just to components handling the current request (controller, view, model)
- **Session scope** data is available across multiple related requests
 - Login to a web mail server
 - e-commerce website with default session until inactivity
 - Session management - implemented as middleware function in module **express-session**

Session

- Session is a mechanism to associate a series of requests coming from a client
 - A conversational state between the client and the server
- HTTP is stateless
 - By default, each request is a session!
 - To maintain a conversational state
 - A server needs to remember what has been going on for EACH client
 - A client needs to send some data to identify self
 - Also a mechanism to control the start/end of a session

How does session work in general

- Client sends the first request
- The server creates an ID for the client and a session object to store session data
 - A server can maintain many sessions simultaneously (identified by ID)
- The server executes predefined business logic for that request and sends back the response together with the ID to the client
- The client stores the ID and associates it with the server
 - A client may maintain many sessions with different servers simultaneously, hence it is important to remember which ID belongs to which server
- When the client sends a second request to this server, it attaches the ID with the request
- The server extracts the ID and use it to find the session data associated with this particular client and make it available to the current request

On the server side

- How can server remembers client state?
 - An object to hold conversational state across multiple requests from the same client identified by a key or ID.
 - It stays for an entire session with a specific client
 - We can use it to store everything about a particular client.
 - Stay in memory mostly but can be persisted in a database

Where does clients stores the ID?

- A **cookie** is a small piece of information stored on a client's computer by the **browser**
- Each browser has its own way to store cookies either in a **text file** or in a **lightweight database**
- Each browser manages its own cookies.
- Since a browser stores cookies from various websites, it also needs a way to identify cookie for a particular site.
- Cookies are identified by {name, domain, path}
- The following are two cookies from different domains

cookie 1

name = **connect.id**

value = s%3AKTObttJqW0k6aVrHB

domain = **localhost**

path = **/**

cookie 2

name = **name**

value = **Joe**

domain = **web.it.usyd.edu.au**

path = **/~comp5347/doc/**

Associate web sites/pages and Cookies

- Browser would associate/send all cookies in the URL scope:
 - cookie-domain is domain-suffix of URL-domain, and
 - cookie-path is prefix of URL-path
- An example
 - Page <http://web.cs.usyd.edu.au/~comp5347/doc/cookie.html> will have cookie B and C, cookie A is not associated with this page

cookie A

name = **name**

value = **Paul**

domain = **web.it.usyd.edu.au**

path = **/~info5010/comp5347/**

cookie B

name = **name**

value = **Joe**

domain = **web.cs.usyd.edu.au**

path = **/~comp5347/doc/**

cookie C

name = **_utma**

value = **223117855...**

domain = **usyd.edu.au**

path = **/**

express-session

- Express application's session management is implemented as middleware function in module **express-session**
- Module **express-session** uses cookie based session management where a small cookie is created to store session id.

Session-aware survey

- Requirements: to ensure that a user cannot vote more than once in a certain period of time ‘using the same browser’.
- Simple solution:
 - Add a session scope variable **vote** to store a user’s previous vote
 - Submit button → check if variable **vote** exists, if true, discard the current vote; else, set the variable to the current vote in session scope and update the results
 - Session scope (object) is accessible to all request as: **req.session**

Session-aware survey

surveysession.server.controller.js

```
var express = require('express')

module.exports.showForm=function(req,res){
  products = req.app.locals.products
  res.render('surveysession.pug',{products:products})
}

module.exports.showResult=function(req,res){
  gender = req.body.gender
  productidx = req.body.vote;
  products = req.app.locals.products;
  surveyresults = req.app.locals.surveyresults;
  sess=req.session;
  if ("vote" in sess)
    res.render('surveysessionresult.pug', {products: products, surveyresults:
surveyresults})
  else{
    sess.vote = productidx;
    gender = req.body.gender
    productidx = req.body.vote;
    if (gender == 0)
      surveyresults.mp[productidx]++;
    else
      surveyresults.fp[productidx]++;
    res.render('surveysessionresult.pug', {products: products, surveyresults:
surveyresults})
  }
}
```

Routes and server.js

surveysession.server.routes.js

```
var express = require('express')
var router = express.Router()
var controller = require('../controllers/surveysession.server.controller')

router.get('/', controller.showForm)
router.post('/survey', controller.showResult)
module.exports = router
```

server.js

```
var express = require('express');
var path = require('path');
var bodyParser = require('body-parser');
var session = require('express-session');

var surveysession = require('./routes/surveysession.server.routes')

var app = express()
app.locals.products=['iphone 7', 'huawei p9', 'Pixel XL', 'Samsung S7']
app.locals.surveyresults = {fp:[0,0,0,0],mp:[0,0,0,0]}

app.set('views', path.join(__dirname, 'views'));
app.use(express.static(path.join(__dirname, 'public')));
app.use(bodyParser.json());
app.use(bodyParser.urlencoded());
app.use(session({secret: 'ssshh', cookie:{maxAge:60000}}));
app.use('/session', surveysession)
app.listen(3000, function () {
  console.log('survey app listening on port 3000!')
})
```

The session will expire in
10 minutes

Cookies sent by server

× Headers Preview Response Cookies Timing

▼ General

Request URL: http://localhost:3000/session/survey

Request Method: POST

Status Code: 🟢 200 OK

Remote Address: [::1]:3000

Referrer Policy: no-referrer-when-downgrade

▼ Response Headers

[view source](#)

Connection: keep-alive

Content-Length: 528

Content-Type: text/html; charset=utf-8

Date: Thu, 27 Apr 2017 05:28:28 GMT

ETag: W/"210-qvEhtChv4kCPeGUXPKpf8P04/8"

set-cookie: connect.sid=s%3AcSdR88-T6fheXi0LfxJSjDkI_z86X9LQ.h0xUAabhDZqUf8irgdnntTTlzb1%2B6aFWjyVmDZxeiY8; Path=/; Expires=Thu, 27 Apr 2017 05:38:28 GMT; HttpOnly

X-Powered-By: Express

Cookies and site data

Site

Locally stored data

[Remove all shown](#)

localhost

localhost

1 cookie

connect.sid

Name:	connect.sid
Content:	s%3AcSdR88-T6fheXi0LfxJSjDkI_z86X9LQ.h0xUAabhDZqUf8irgdnntTTlzb1%2B6aFWjyVmDZxeiY8
Domain:	localhost
Path:	/
Send for:	Any kind of connection
Accessible to script:	No (HttpOnly)
Created:	Thursday, April 27, 2017 at 3:28:28 PM
Expires:	Thursday, April 27, 2017 at 3:38:28 PM

[Remove](#)

▼ Request Headers

[view source](#)

Accept: text/html,application/xhtml+xml,application/xml;q=

Accept-Encoding: gzip, deflate, br

Accept-Language: en-US,en;q=0.8

Cache-Control: no-cache

Connection: keep-alive

Content-Length: 15

Content-Type: application/x-www-form-urlencoded

Cookie: connect.sid=s%3AKT0bttJqW0k6aVrHBQ80po9ZPY0qQbG7.6

Outline

- Implementing MVC basics
 - Application Folder Structure
 - CommonsJS modules
 - Controller and Routers
- Session Management
- Database Layer
 - Introduction to NoSQL
 - Introduction to MongoDB

NoSQL Brief Introduction

- NoSQL encompasses the general trend of a new generation of database servers
 - Challenges of traditional SQL databases
 - Scalability, flexible schema, object relational mismatch, etc
- Broad categories of NoSQL systems
 - Document Storage (e.g. MongoDB)
 - Key-Value Storage
 - Column based Storage
 - Graph database

Document Storage and MongoDB

- Document storage system based on self describing documents
 - Entity stored as a document (record (row) in typical SQL0
 - Two dominant formats
 - XML
 - JSON (JavaScript Object Notation)

```
Invoice _1= {  customer: {name: "John", address: "Sydney"},  
              product: { code: "123", quantity: 2}  
            }
```

```
Invoice _2= {  customer: {name: "Smith", address: "Melbourne"},  
              product: { code: "xyz", quantity: 20},  
              delivery: "express"  
            }
```

JSON Data Format

- JSON (**J**ava**S**cript **O**bject **N**otation): represent JavaScript objects as strings
- Introduced in 1999 as an alternative to XML for data exchange
- JSON object: a list of property names and values contained in curly braces:
`{ propertyName1 : value1, propertyName2 : value2 }`
- Arrays represented with square brackets
`[value1, value2, value3]`

Matching Terms in SQL and MongoDB

SQL	MongoDB
Database	Database
Table	Collection
Index	Index
Row	BSON* document
Column	BSON field
Primary key	_id field
Join	Embedding and referencing \$lookup in aggregation (since 3.2)

*Binary JSON

MongoDB Document Model – Example

<u>TFN</u>	Name	Email	age
12345	Joe Smith	joe@gmail.com	30
54321	Mary Sharp	mary@gmail.com	27

```
{ _id: 12345,  
  name: "Joe Smith",  
  email: "joe@gmail.com",  
  age: 30  
}
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27  
}
```

MongoDB Document Model - Example

users table in RDBMS

Column name is part of schema

TFN	Name	Email	age
12345	Joe Smith	joe@gmail.com	30
54321	Mary Sharp	mary@gmail.com	27

two rows

{ _id : 12345, name: "Joe Smith", email: "joe@gmail.com", age: 30 }
{ _id : 54321, name: "Mary Sharp", email: "mary@gmail.com", age: 27 }

two documents

Field name is part of data

Native Support for Array

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30  
}
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27  
}
```

<u>TFN</u>	Name	Email	age
12345	Joe Smith	joe@gmail.com , joe@ibm.com ??	30
54321	Mary Sharp	mary@gmail.com	27

Native Support for Embedded Document

```
{_id: 12345,  
  name: "Joe Smith",  
  email: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30  
}
```

```
{_id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27,  
  address: { number: 1,  
             name: "cleveland street",  
             suburb: "chippendale",  
             zip: 2008  
          }  
}
```

<u>TFN</u>	Name	Email	age
12345	Joe Smith	joe@gmail.com	30
54321	Mary Sharp	mary@gmail.com	27

Native Support for Embedded Document

```
{_id: 12345,  
  name: "Joe Smith",  
  email: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30  
}
```

```
{_id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27,  
  address: { number: 1,  
             name: "cleveland street",  
             suburb: "chippendale",  
             zip: 2008  
          }  
}
```

<u>TFN</u>	Name	Email	age	address
12345	Joe Smith	joe@gmail.com	30	
54321	Mary Sharp	mary@gmail.com	27	1 cleveland street, chippendale, NSW 2008

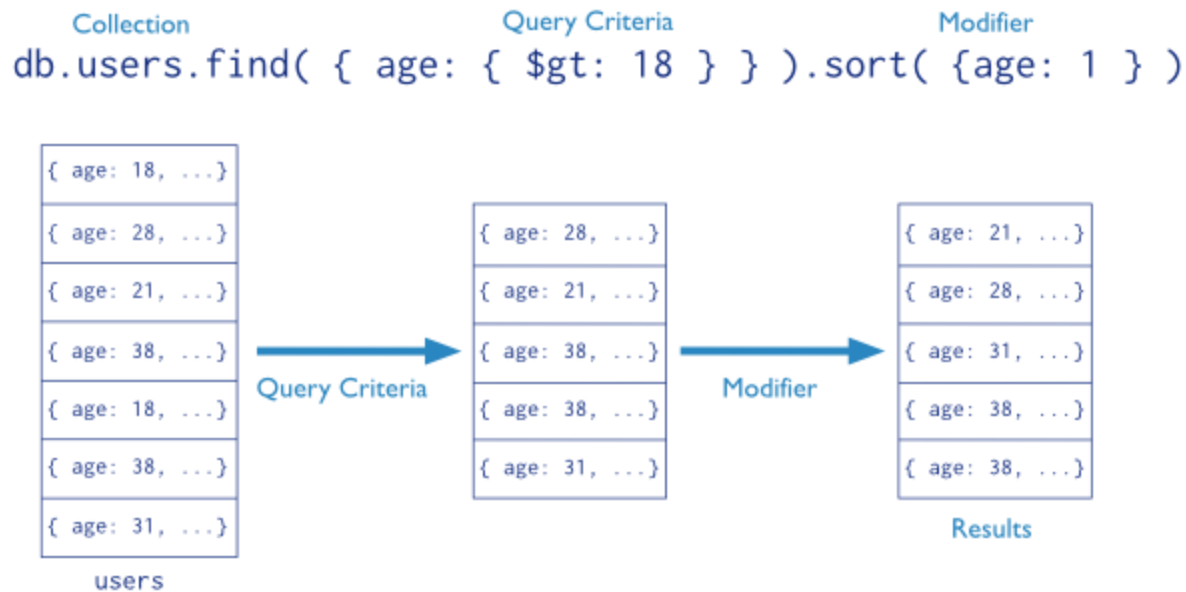
MongoDB data types

- Primitive types
 - String, integer, boolean (true/false), double, null
- Predefined special types
 - Date, object id, binary data, regular expression, timestamp,
 - DB Drivers implement them in language-specific way
 - The interactive shell provides constructors for all
- Array and object
- Field name is of **string** type with certain restrictions
 - “_id” is reserved for primary key
 - cannot start with “\$”, cannot contain “.” or null

MongoDB Queries

- A ***read*** query targets a specific collection
 - **criteria**
 - may include a **projection** to specify fields from the matching documents
 - may include **modifier** to limit, skip, or sort the results
- A ***write*** query may *create, update or delete* data
 - One query modifies the data of a single collection
 - Update and delete query can specify query **criteria**

Read Query Example



Find documents in the **users** collection with **age** field greater than 18, sort the results in ascending order by **age**

Read Query Interface

– `db.collection.find()`

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

Read Query Interface

– db.collection.find()

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

Find at most 5 documents in the **users** collection with **age** field greater than 18, return only the name and address field of each document.

Read Query Interface

– db.collection.find()

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

Find at most 5 documents in the **users** collection with **age** field greater than 18, return only the name and address field of each document.

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

← projection
← table
← select criteria
← cursor modifier

Read Query Features

- Find data using any criteria
 - Does not require indexing
 - Indexing can improve performance
 - JOIN is not supported!!
- Query **criteria** expressed as BSON document (query object)
 - Individual expressed using predefined selection operator
 - Eg. `$lt` is the operator for “less than”
- Query **projection** are expressed as BSON document

Read Query Features

SQL	MongoDB Query in Shell
<code>select * from user</code>	<code>db.user.find()</code> or <code>db.user.find({})</code>
<code>select name, age from user</code>	<code>db.user.find({}, {name: 1, age: 1, _id: 0})</code>
<code>select * from user where name = "Joe Smith"</code>	<code>db.user.find({name: "Joe Smith"})</code>
<code>select * from user where age < 30</code>	<code>db.user.find({age: {\$lt: 30}})</code>

Querying Array field

- Like querying simple type field
 - `db.user.find({emails: "joe@gmail.com"})`
 - `db.user.find({"emails.0": "joe@gmail.com"})`

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30}
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27}
```

<http://docs.mongodb.org/manual/tutorial/query-documents/#arrays>

Querying Embedded Document

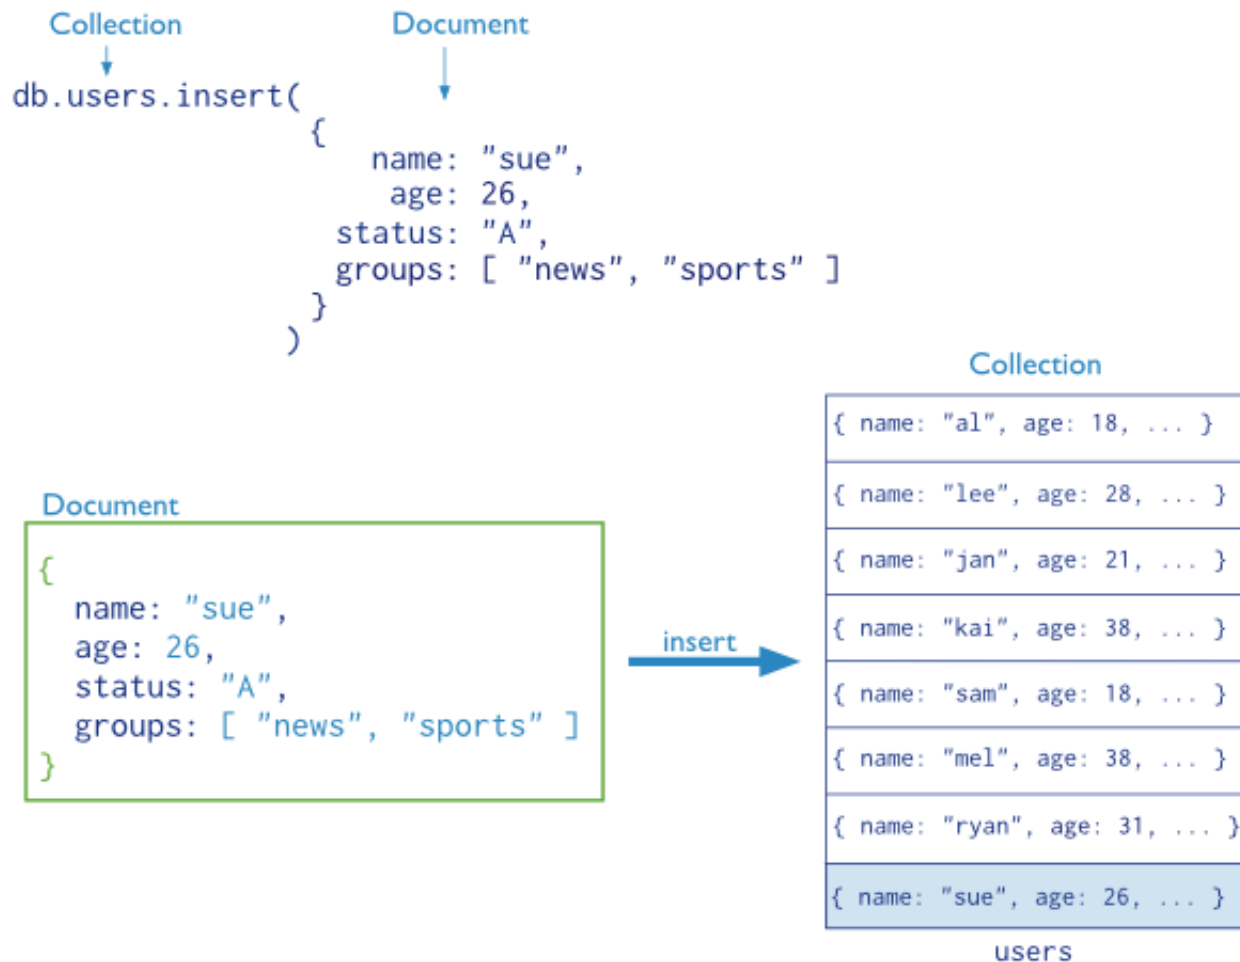
- Queried as a **whole**, by **individual field**, or by **combination of individual fields**
 - `db.user.find({address: {number: 1, name: "pine street", suburb: "chippendale", zip: 2008}})`
 - `db.user.find({"address.suburb": "chippendale"})`
 - `db.user.find({address: {$elemMatch: {name: "pine street", suburb: "chippendale"}}})`

```
{ _id: 12345,  
  name: "Joe Smith", email: ["joe@gmail.com", "joe@ibm.com"], age: 30,  
  address: {number: 1, name: "pine street", suburb: "chippendale", zip: 2008 }  
}
```

```
{ _id: 54321,  
  name: "Mary Sharp", email: "mary@gmail.com", age: 27,  
  address: { number: 1, name: "cleveland street", suburb: "chippendale", zip: 2008 }  
}
```

<http://docs.mongodb.org/manual/tutorial/query-documents/#embedded-documents>

Write Query- Insert



Insert a new document in **users** collection.

Insert Example

- `db.user.insert({_id: 12345, name: "Joe Smith", emails: ["joe@gmail.com", "joe@ibm.com"], age: 30})`
- `db.user.insert({ _id: 54321, name: "Mary Sharp", email: "mary@gmail.com", age: 27, address: { number: 1, name: "cleveland street", suburb: "chippendale", zip: 2008}})`

Insert Example

- `db.user.insert({_id: 12345, name: "Joe Smith", emails: ["joe@gmail.com", "joe@ibm.com"], age: 30})`
- `db.user.insert({_id: 54321, name: "Mary Sharp", email: "mary@gmail.com", age: 27, address: { number: 1, name: "cleveland street", suburb: "chippendale", zip: 2008}})`

user collection

```
{ _id: 12345, name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30  
}  
  
{ _id: 54321,  
  name: "Mary Sharp", email: "mary@gmail.com", age: 27,  
  address: { number: 1,  
             name: "cleveland street",  
             suburb: "chippendale",  
             zip: 2008  
            }  
}
```

Insert Behaviour

- If the new document does not contain an “_id” field, the system will add an “_id” field and assign a unique value to it
- If the new document does contain an “_id” field, it should have a unique value

Write Operation - Update

```
db.users.update(  
  { age: { $gt: 18 } },  
  { $set: { status: "A" } },  
  { multi: true }  
)
```

← collection
← update criteria
← update action
← update option

Has the same effect as the following SQL:

```
UPDATE users  
SET    status = 'A'  
WHERE  age > 18
```

← table
← update action
← update criteria

Updates operators

- Modifying simple field: \$set, \$unset
 - `db.user.update({_id: 12345}, {$set: {age: 29}})`
 - `db.user.update({_id: 54321}, {$unset: {email: 1}}) // remove the field`

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 29}
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27}
```

<http://www.mongodb.org/display/DOCS/Updating>

Updates operators

- Modifying array elements: `$push`, `$pushAll`, `$pull`, `$pullAll`
 - `db.user.update({_id: 12345}, {$push: {emails: "joe@hotmail.com"}})`
 - `db.user.update({_id: 54321},
{$pushAll: {emails: ["mary@gmail.com", "mary@microsoft.com"]}})`
 - `db.user.update({_id: 12345}, {$pull: {emails: "joe@ibm.com"}})`

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30 }
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27 }
```

Updates operators

- Modifying simple field: \$set, \$unset
 - db.user.update({_id: 12345}, {\$set: {age: 29}})
 - db.user.update({_id: 54321}, {\$unset: {email: 1}}) // remove the field
- Modifying array elements: \$push, \$pushAll, \$pull, \$pullAll
 - db.user.update({_id: 12345}, {\$push: {emails: "joe@hotmail.com"}})
 - db.user.update({_id: 54321},
{\$pushAll: {emails: ["mary@gmail.com", "mary@microsoft.com"]}})
 - db.user.update({_id: 12345}, {\$pull: {emails: "joe@ibm.com"}})

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@ibm.com"],  
  age: 30 }
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  email: "mary@gmail.com",  
  age: 27 }
```

```
{ _id: 12345,  
  name: "Joe Smith",  
  emails: ["joe@gmail.com", "joe@hotmail.com"],  
  age: 29 }
```

```
{ _id: 54321,  
  name: "Mary Sharp",  
  emails: ["mary@gmail.com", "mary@microsoft.com"],  
  age: 27 }
```

<http://www.mongodb.org/display/DOCS/Updating>

Write Operation - Delete

- `db.user.remove();`
 - Remove all documents in user collection
- `db.user.remove({_id: 12345})`
 - Remove document with a particular id from user collection

Aggregation

- Simple and relatively standard data analytics can be achieved by **aggregation**
 - Grouping, summing up value, counting, sorting, etc
 - Running on the DB engine instead of application layer
- Several options
 - Aggregation Pipeline
 - MapReduce
 - JavaScript functions
 - Performance
 - Customized aggregations

Aggregation Pipeline

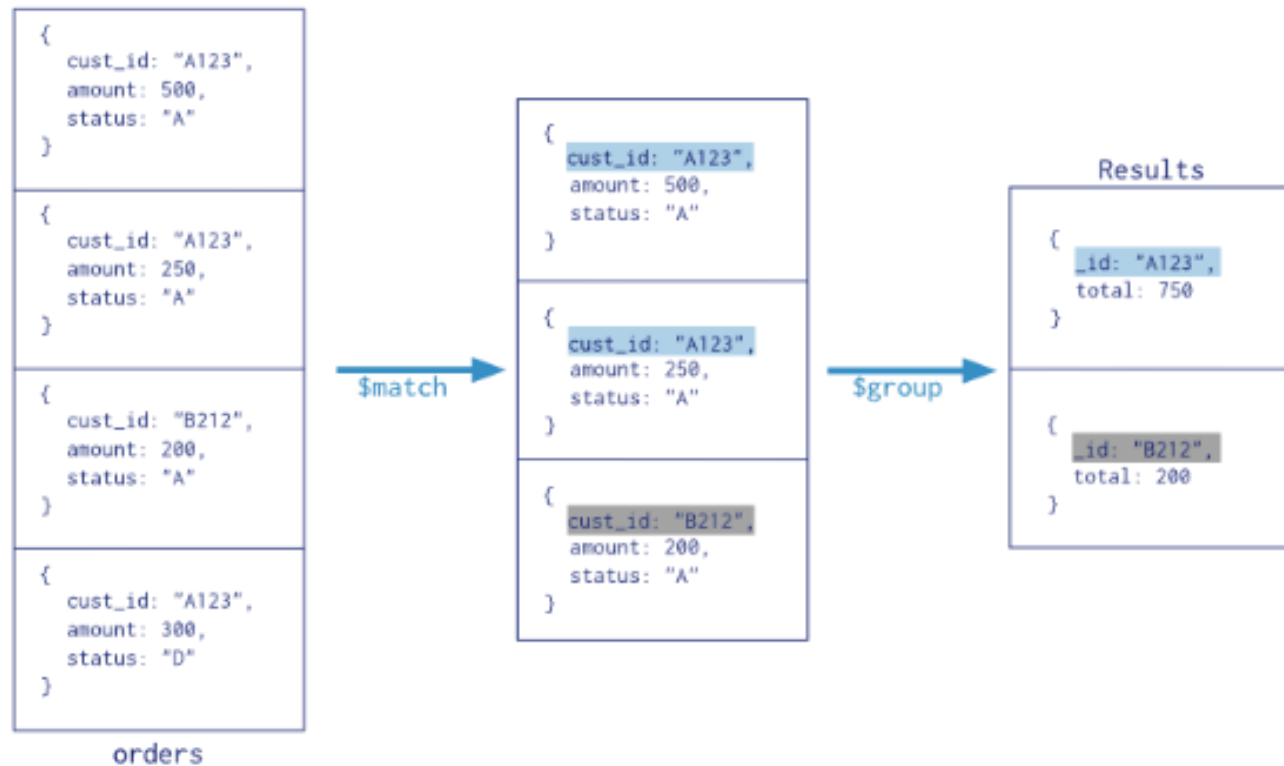
- Consists of multiple stages
 - Specified using **pipeline operators** (e.g., **\$match**, **\$group**, **\$sort** and so on)
 - Similar to SQL's WHERE, GROUP BY, SORT BY etc
 - Each stage is expressed as an object enclosed by curly bracket
 - Various **expressions** can be specified in each stage
 - To filter documents or to perform simple calculation on a document
 - **\$substr**, **\$size**, etc, ...
 - **\$group** stage can specify **accumulators** to perform calculation on documents with the same group key

Aggregation Pipeline – Format

```
db.collection.aggregate( [  
    { pipeline operator: {expression/accumulator,..., expression/accumulator} },  
    { pipeline operator: {expression/accumulator,..., expression/accumulator} },  
    ...  
] )
```

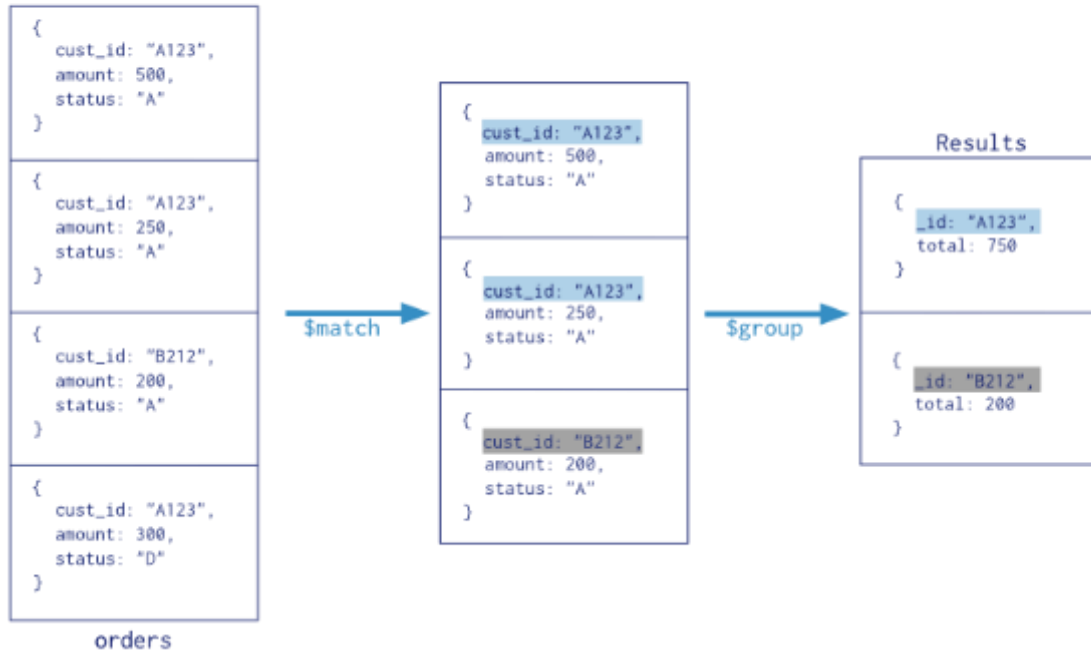

Aggregation Example

Collection
↓
`db.orders.aggregate([`
 `$match stage → { $match: { status: "A" } },`
 `$group stage → { $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]`)



Aggregation Example

Collection
↓
`db.orders.aggregate([`
 \$match stage → `{ $match: { status: "A" } },`
 \$group stage → `{ $group: { _id: "$cust_id", total: { $sum: "$amount" } } }`
 `]`)



```
select cust_id as _id, SUM(amount) as total
from orders
where status = "A"
group by cust_id
```

Aggregation Behaviour

- Operates on a single collection (before 3.2)
 - Join can be performed using a particular operator **\$lookup**
- Logically passes the entire collection into the pipeline
- Early filtering can improve the performance
- **\$match** and **\$sort** operators are able to use index if placed at the beginning of the pipeline

Resources

- Haviv, Amos Q, MEAN Web Development
 - E-book, accessible from USYD library
 - Chapter 4 and 5
- MongoDB online documents:
 - MongoDB CRUD Operations
 - <http://docs.mongodb.org/manual/core/crud-introduction/>
 - MongoDB Aggregation
 - <http://docs.mongodb.org/manual/core/aggregation-introduction/>

Week 6 Tutorial: Node.js MVC Application

Week 7 Lecture: Connecting to MongoDB

